

An Introduction to Differential Cryptanalysis of Block Ciphers

Steven Rybicki (RYBSTE001)

Abstract

Block ciphers, ciphers that take blocks of input, are used in cryptography to encrypt data and build other cryptographic schemes. This is only possible if block ciphers have the confidentiality property: any text encrypted under a secret key with that block cipher is not able to be decrypted without that key. In this project we discuss differential cryptanalysis and how this can be used to break the confidentiality of certain substitution-permutation network based block ciphers. We start by outlining the basic results in probability and cryptanalysis of substitution-permutation networks required to understand the attack. We then go through the details of how to attack a cipher using differential cryptanalysis, building the intuition required to understand how this attack works and then applying this to a toy cipher.

Contents

1	Preface	4
2	Discrete Probability Theory	5
2.1	Introduction	5
2.2	Random Variables	8
2.3	Conditional Probability	8
2.4	Independence	9
3	Introduction to Cryptanalysis	11
3.1	Block Ciphers	11
3.2	Confidentiality, OTP and Xor	14
3.3	Chosen-plaintext attack and Kerckhoffs's principle	16
4	Basic Elements of Differential Cryptanalysis	18
4.1	Outline of an Attack using Differential Cryptanalysis	18
4.2	Differential Cryptanalysis on One Round	19
4.2.1	Key Addition	19
4.2.2	Substitution	20
4.2.3	Permutation	22
4.2.4	Characteristics	23
4.3	Multiple Rounds	24
4.3.1	S-round Characteristics	24
4.3.2	Markov Chains and Ciphers	25
4.3.3	Differentials	29
5	Preparing for an Attack	32
5.1	Finding Good Differentials	32
5.2	Filtering	34
6	Attacking a Cipher using Differential Cryptanalysis	37
6.1	Generic One Round Attack	37
6.1.1	Algorithm	37

6.1.2	Explanation	38
6.1.3	Extensions of Generic Attack	39
6.2	Attack on Unsub A	39
6.2.1	Data Complexity and chance for Success	41
A	Unsub A Construction	45
A.1	Cipher Description	45
A.2	Round Structure	45
A.2.1	Key Addition	45
A.2.2	Substitution	45
A.2.3	Permutation	45
A.2.4	Key Schedule	46
A.3	Encryption and Decryption Algorithms	46

List of Figures

1	Visualisation of UnsubA Permutation emphasising Bit Location	52
2	Visualisation of UnsubA Permutation emphasising Sbox Origin	53
3	Visualisation of UnsubA Encryption	54

List of Tables

1	Substitution Layer Definition for UnsubA	49
2	Permutation Layer Definition for UnsubA	50
3	Substitution Layer Example Values for UnsubA	50
4	Permutation Layer Example Values for UnsubA	51

1 Preface

Cryptanalysis is the study of recovering information from ciphers without knowing the secret key. *Differential cryptanalysis* tries to recover secret information from ciphers by looking at how the “differences” in pairs of inputs can affect the differences of the outputs. This field, while first publicly disclosed by Biham and Shamir [1], was identified earlier and kept secret by the designers of DES (Data Encryption Standard) as it was regarded as powerful enough of a weapon against ciphers to be a national secret. Since then it has been used to great effect against a number of ciphers.

While you can perform differential cryptanalysis on various kinds of ciphers, in this project we are going to focus on applying it to block ciphers, specifically iterated block ciphers designed as a substitution-permutation network like Unsub A (see Appendix A). These terms are defined later in the project.

In order to understand these attacks fully, one needs some background in discrete probability and cryptanalysis. For those readers who have not have been exposed to these concepts, we have provided a primer to get you up to speed in Sections 2 and 3. If you are familiar with basic cryptanalysis concepts and discrete probability, or are not interested in the theoretical basis for the attacks, simply skim over these sections and start reading in earnest starting with Section 4.

In each section relating to the differential cryptanalysis and attack, we will first try to develop the theoretical basis that can applied to all block ciphers based on a substitution-permutation network. Where relevant we will also give concrete examples of that attack or analysis as performed on Unsub A. This is to ensure the reader has an understanding of why the analysis or attack is relevant and valid as well as an example to ground this understanding.

We are going to be basing our approach on the The Block Cipher Companion [4]. All other sources will be cited when mentioned.

2 Discrete Probability Theory

2.1 Introduction

We all have an intuitive sense of probability which we can use to predict the likelihood of events: we know that winning the lottery is unlikely and that a fair coin when flipped should come up heads as often as it comes up tails. In each of these, we have an event (winning the lottery, a coin coming up heads) and an associated probability (a small one and $\frac{1}{2}$). In this section, we are going to take our intuitive understanding of probability and formalize it.

Since this project is not a text on probability, we will give only a brief introduction. If you want a more comprehensive treatment, see any good mathematical text on discrete probability theory such as *Grinstead and Snell's Introduction to Probability* [3]. This text is available for free under the GNU Free Document Licence ¹.

Let us first look at what it means to say that “An event X will occur with Probability Y ”.

Definition 1. We call any procedure that a) we can repeat and b) has a set of well defined outcomes an **experiment**. For any experiment, let the **sample space** denote the all outcomes of this experiment. We will usually represent the sample space by Ω . If X is a subset of Ω , we call X an **event**.

Let us take two every day events, flipping a fair coin and rolling a die, and represent them using the concepts we introduced above.

Example 2. In the experiment of “flipping a fair coin”, the sample space would be the coin coming up heads or tails². If we represent the result of a coin coming up heads by H and similarly represent the result of a coin landing on tails by T , we have that for this experiment $\Omega = \{H, T\}$. Here, our events are $\{H\}, \{T\}$ (the coin coming up heads or tails respectively), \emptyset (the coin comes up neither heads or tails) or Ω (the coin comes up heads or tails).

Note that even though \emptyset will never occur, it is still a valid event as $\emptyset \subset \Omega$ for any sample space Ω .

Example 3. If we looked at the experiment of rolling a traditional six sided die and examining its displayed value, we would have $\Omega = 1, 2, 3, 4, 5, 6$ with the event $\{2, 4, 6\}$ corresponding to rolling an even number.

Now, with both of the above examples, each outcome occurs with equal likelihood. Rolling a 6 on a die is as likely as rolling a 1; flipping a fair coin and getting heads is

¹At the time of writing, a copy is available at the following link: http://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/book.html

²This holds if we ignore a coin landing on its side which in this case we will.

as likely as getting tails. However, not all elements in a sample space occur with equal likelihood as our example below shows.

Example 4. *Let us examine the experiment of you and a friend each picking an integer from one to 10^{10} at random and seeing if they were equal. We would have*

$$\Omega = \{\text{Numbers are equal}, \text{Numbers are not equal}\}$$

In this sample space, the event “{Numbers are not equal}” is more likely to occur than the event {Numbers are equal}.

For the dice and fair coins, we say that there is a *uniform probability distribution* - all individual outcomes occur with equal likelihood. However, this is not always the case. So before we can define what a probability is, we need to look at *probability distributions*.

Definition 5 (Probability Distribution). *A **probability distribution** for an experiment with sample space Ω is a function $m : \Omega \rightarrow [0, 1]$ such that*

$$\sum_{x \in \Omega} m(x) = 1$$

*If $(\forall x \in \Omega) m(x) = \frac{1}{|\Omega|}$, then m is said to be a **uniform probability distribution**.*

With this in mind, we can define the concept of a probability.

Definition 6 (Probability). *Given an experiment with sample space Ω and probability distribution m over Ω , the **probability** of an event X is defined as a function $Pr : 2^\Omega \rightarrow [0, 1]$ where 2^Ω is the **powerset** of Ω (set of all subsets of Ω) and*

$$Pr[X] = \sum_{x \in X} m(x)$$

Here, the powerset of our sample space Ω is simply the set of all events. So in the case of Example 2, we have that

$$2^\Omega = \{\emptyset, \{H\}, \{T\}, \Omega\}$$

The powerset of X is also sometimes represented as $P(X)$, but we prefer the 2^X notation as the number of elements in the powerset of a finite set is X is $2^{|X|}$.

Loosely, the probability of an event A corresponds to the likelihood that A occurs in our experiment where as our probability distribution corresponds to the likelihood that a singleton from Ω occurs.

Note that for a given probability distribution, the probability for an event is fixed. Therefore if we have a case where we want to change the probability associated with an event, we must change our probability distribution. This is illustrated in the example below.

Example 7. Let us consider flipping a fair coin versus flipping a biased coin (a coin which lands on one side more often than another). Both have the same sample space $\Omega = \{H, T\}$ but the probability of the coin landing heads for each will be different. We represent this by giving them different probability distributions. For the fair coin, we will have a uniform distribution. For the biased coin, we would have $m(H) = p, m(T) = 1 - p$ for some $p \in [0, 1]$.

There is an alternate way to introduce the concept of a probability which involves using *Kolmogorov's Axioms*. We are going to list these as a theorem and prove them using the concept of probability defined above. Note that as we go on, we will stop explicitly mentioning we are talking about experiments and simply refer to their sample spaces, with the experiment being made obvious by the context.

Theorem 8 (Kolmogorov). *Given a sample space Ω for an experiment with a probability distribution m over Ω and events $A, B \subseteq \Omega$, we have that*

1. $0 \leq Pr[A] \leq 1$
2. $Pr[\Omega] = 1, Pr[\emptyset] = 0$
3. If $A \cap B = \emptyset$ then $Pr[A \cup B] = Pr[A] + Pr[B]$

As this is a project on differential cryptography rather than probability, only the sketches for the proofs are provided below.

Proof. We prove each of these separately.

1. Note that $Pr[X] = \sum_{x \in X} m(x)$. Since m has a codomain of $[0, 1]$, for any $x \in \Omega, m(x) \geq 0$ which means that $\sum_{x \in X} m(x) \geq 0$. For the other inequality, we note that

$$\sum_{x \in X} m(x) \leq \sum_{x \in X} m(x) + \sum_{x \in (\Omega - X)} m(x) = \sum_{x \in \Omega} m(x) = 1$$

2. $Pr[\Omega] = \sum_{x \in \Omega} m(x) = 1$ and $Pr[\emptyset] = \sum_{x \in \emptyset} m(x) = 0$ as there are no elements in the empty set.
3. We will prove this for finite Ω . Let us set $A = \{a_0, \dots, a_n\}$ and $B = \{b_0, \dots, b_m\}$. Then we have that

$$Pr[A \cup B] = \sum_{x \in A \cup B} m(x) = m(a_0) + \dots + m(a_n) + m(b_0) + \dots + m(b_m) = \sum_{x \in A} m(x) + \sum_{x \in B} m(x)$$

as their intersections are empty. This gives $Pr[A \cup B] = Pr[A] + Pr[B]$.

□

2.2 Random Variables

In the previous section, we developed a language which we can use to describe the probability of a specific event happening according to a probability distribution over a sample space. However, we would like a more flexible way to specify events like events that occur randomly.

For example, consider the experiment of taking a card out of a deck. How would you express the event “Picked a red card”? We would simply create a set which contain all the red cards. However, this approach becomes a bit less elegant when we look at events like “Picked a number card with value i ” for some i in $\{1, 2, \dots, 9, 10\}$. It is possible to do so with the concepts developed above, but we would prefer a much more elegant way. To do this, we use *Random Variables*.

Definition 9 (Random Variables). *For a sample space Ω , a **random variable** is a function $r : \Omega \rightarrow V$ where V is a non-empty set.*

For $v \in V$ we define the probability of the random variable r taking the value v over Ω to be $Pr[r = v] = Pr[r^{-1}(v)]$.

*If $V = \Omega$ and r is Id_Ω , i.e. the identity function on Ω , then r is said to be the **uniform random variable** over Ω .*

Example 10. *Let us take the card experiment from above. To describe the event of picking a red card, we would define a random variable r to go to the set $\{1, 0\}$ with r being defined as sending all red cards to 1 and all black cards to 0. Then, we can define the probability desired as $Pr[r = 1]$. Similarly, if we wanted to define the event “Picked a number card with value 2”, we would set $V = \{0, 2, \dots, 9, 10\}$ with r sending all number cards to their values and all face cards (Ace, King, Queen, Jack) to 0. Then the probability for the event would be*

$$Pr[r = 2] = Pr[r^{-1}(2)] = Pr[\{2 \text{ of clubs}, 2 \text{ of spades}, 2 \text{ of hearts}, 2 \text{ of diamonds}\}] = \frac{1}{13}$$

as the probability distribution for picking cards is uniform.

2.3 Conditional Probability

We now have a flexible way to describe probability distributions, but still no elegant way to describe a sequence of events. For example, we might want to know what the probability is that we draw an ace from a pack of cards given that in the previous three draws we drew all aces. In order to do this, we need to define what it means for something to follow conditionally from something else. We also need some way of calculating this probability.

Definition 11 (Conditional Probability). *Let A and B be events in Ω . Then we represent the probability of A occurring after B is known to have occurred as*

$$Pr[A|B] = \frac{Pr[A \cap B]}{Pr[B]}$$

This is known as **conditional probability**.

To give some intuition for this definition, let us reason about the quantity we are calculating. Since we are looking at the event that A occurs after B has already occurred, it makes sense that $P[A|B]$ be proportional to $Pr[A \cap B]$, the probability that both A and B occur. However, $Pr[A \cap B]$ under represents this probability. To see this, consider the case where $Pr[B] = a$ and $Pr[A] = 1$. Then $Pr[A|B]$ should be 1 as the probability that A occurs is 1 regardless of whether B occurs before it or not. However, $Pr[A \cap B] = a$, an underrepresentation of the probability desired by a factor of $a = Pr[B]$.

Note that here we are using \cap as the set intersection operator as A, B are just subsets of Ω . Sometimes instead of using \cap , we will abbreviate this and list the sets separated by commas like so:

$$Pr[A \cap B | C \cap D] = Pr[A, B | C, D]$$

While this gives us one way of calculating a conditional probability, another that is quite useful is derived in what is often called *Bayes' Theorem*.

Theorem 12 (Bayes'). *Given events A and B in Ω , we have that*

$$Pr[A|B] = \frac{Pr[B|A]Pr[A]}{Pr[B]}$$

Proof.

$$Pr[A|B] = \frac{Pr[A \cap B]}{Pr[B]} = \frac{Pr[A \cap B]Pr[A]}{Pr[B]Pr[A]} = \frac{Pr[A \cap B]}{Pr[A]} \frac{Pr[A]}{Pr[B]} = \frac{Pr[B|A]Pr[A]}{Pr[B]}$$

An alternate way to see this is to note that

$$Pr[A \cap B] = Pr[B]Pr[A|B] = Pr[A]Pr[B|A]$$

If we multiple second two equalities by $Pr[B]^{-1}$ then the result follows. \square

2.4 Independence

Let us say some bored statistician decided to flip a fair coin one hundred times. To her surprise, the first ninety nine coin tosses came up tails. What is the chance that her hundredth throw is going to come up tails?

Here, we are not asking what is the chance of flipping a hundred fair coins and having one hundred of them coming up tails. We are asking for the hundredth flip, what is the

chance of that coin coming up heads given the previous ninety nine have come up tails? If we call the event of the coin coming up tails A and the coin coming up tails ninety nine times B what we are asking is what the value of the following expression is:

$$Pr[A|B] = \frac{Pr[A \cap B]}{Pr[B]}$$

Now, we know that the chance of a fair coin coming up heads or tails is always $\frac{1}{2}$. This is the case whether there is been a hundred or a thousand flips before it. Therefore we can tell the statistician that despite her previous throws, her last throw coming up tails will occur with probability $\frac{1}{2}$. In terms of our definition above, this is saying that $Pr[A|B] = Pr[A]$.

Put in another way, what we are asserting is that A is *independent* of B - that the probability for A does not depend on B in any way. This is captured in the following definition.

Definition 13 (Independence). *Events A and B are said to be **independent** on the probability distribution m over Ω if*

$$Pr[A \cap B] = Pr[A] \times Pr[B]$$

If we say that our events A and B are independent, then from the definition of conditional probability we have that

$$Pr[A|B] = \frac{Pr[A \cap B]}{Pr[B]} = \frac{Pr[A]Pr[B]}{Pr[B]} = Pr[A]$$

which is exactly the behaviour we wanted. We can also extend this definition to random variables.

Definition 14 (Independence of Random Variables). *If we have two random variables X, Y which take values from the sample space Ω to V then X, Y are **independent** if*

$$(\forall a, b \in V) Pr[X = a, Y = b] = Pr[X = a] \times Pr[Y = b]$$

3 Introduction to Cryptanalysis

In order to give our differential cryptanalysis some context we will be covering some of the basics of block cipher cryptanalysis in this Section. We will start by giving a little background on block ciphers including some ideas on different security models, eventually introducing the attack parameters for our differential cryptanalysis based attack.

This is intended for those completely new to cryptanalysis in order to motivate why some of the concepts we introduce and use are necessary, relevant and realistic. If you understand binary xors, know what a chosen plaintext attack is and why we might want to use that definition of confidentiality over perfect secrecy, then you can safely skim this section. If not, you are encouraged to read on in order to fully understand the differential cryptanalysis in later sections.

3.1 Block Ciphers

Block Ciphers are ciphers that operate on blocks of data of some fixed size. We typically represent these blocks as either *bit strings* - sequences of 1's and 0's - of a certain length n or numbers under 2^n .

Definition 15 (Bit-String). A **bit-string** of length n is a sequence $(x_i)_{i=0}^{n-1}$ where each element of the sequence is either a 1 or a 0. We will denote these bit strings as $0bx$ where x is the string.

Example 16. 0110 is a bit-string of length 4 which we will represent here by 0b0110.

If a number is smaller than $2^n - 1$ then we can represent it as a bit string of length n by taking its representation in base 2 (i.e. its binary representation) and then preappending 0's until it has length n .

Example 17. To represent the number 5 as a length 4 bit string would be 0b0101. We can also convert bit strings to numbers, for example 0b1111 would be 15.

With this, we can see that a bit string is just a representation of a number in a certain range in base 2. This means in this project we will treat the input and output of block ciphers as numbers in certain circumstances and sequences of digits in others, making it clear by the context which form we are working with. We might also represent these numbers in hexademical (numbers in base 16) if they are large. We will represent hexademical numbers as $0xy$ where y is a representation of a bit-string in base 16. An example of counting in hexadecimal, binary and decimal numbers is in the example below.

Example 18. Below we have an example of counting in binary, hexadecimal and decimal to check your understanding on.

Decimal Numbers	Binary Numbers	Hexadecimal Numbers
0	0b0	0x0
1	0b1	0x1
2	0b10	0x2
3	0b11	0x3
4	0b100	0x4
5	0b101	0x5
6	0b110	0x6
7	0b111	0x7
8	0b1000	0x8
9	0b1001	0x9
10	0b1010	0xa
11	0b1011	0xb
12	0b1100	0xc
13	0b1101	0xd
14	0b1110	0xe
15	0b1111	0xf
16	0b10000	0x10
17	0b10001	0x11
18	0b10010	0x12
19	0b10011	0x13
20	0b10100	0x14
21	0b10101	0x15
22	0b10110	0x16
23	0b10111	0x17
24	0b11000	0x18
25	0b11001	0x19
26	0b11010	0x1a
27	0b11011	0x1b
28	0b11100	0x1c
29	0b11101	0x1d
30	0b11110	0x1e
31	0b11111	0x1f
32	0b100000	0x20

Block ciphers take in a block of *message* or *plaintext* m (we will be using these terms interchangeably) and encrypts it into a block of *ciphertext* c using a secret key k . We capture this formally using functions for encryption and decryption.

Definition 19 (Block Cipher). *Let \mathcal{M} be the set containing all valid input bit-strings and let \mathcal{K} be the set of all valid key bit-strings. Then we say that a block-cipher B consists of the pair of encryption and decryption functions $E : \mathcal{M} \times \mathcal{K} \rightarrow \mathcal{M}$ and $D : \mathcal{M} \times \mathcal{K} \rightarrow \mathcal{M}$ such that for $m, c \in \mathcal{M}$ and $k \in \mathcal{K}$ we have that*

$$D(E(m, k), k) = m \text{ and } E(D(m, k), k) = m$$

If we fix some $k \in \mathcal{K}$, then we can define $E_k(m) = E(m, k)$, $D_k(m) = D(m, k)$ for all $m \in \mathcal{M}$ as the encryption and decryption operations of a cipher under a key k .

Technically, one should define these functions' codomains as a separate set \mathcal{C} of all possible ciphertexts. Here we simplify this definition as Unsub A, along with many other block ciphers, have $\mathcal{C} = \mathcal{M}$. You can reformulate all the results here with a separate co-domain set without much difficulty.

UnsubA is known as an iterated block cipher with a substitution permutation network (SPN). This is clarified below.

Definition 20 (Iterated Block Cipher). *An **iterated block cipher** is a block cipher that is composed of multiple applications (rounds) of a round function. An iterated block cipher is a substitution permutation network (SPN) block cipher if its round function consists of a key addition function $K : \mathcal{M} \times \mathcal{K} \rightarrow \mathcal{M}$, a non-linear permutation $S : \mathcal{M} \rightarrow \mathcal{M}$ and a linear permutation $P : \mathcal{M} \rightarrow \mathcal{M}$.*

One interesting way to look at the cipher for a certain key is as permutations on the set of \mathcal{M} . For this to be valid we need E_k, D_k for all $k \in \mathcal{K}$ to be bijective which we prove below.

Lemma 21. *For all $k \in \mathcal{K}$, E_k and D_k are bijective.*

Proof. For $m \in \mathcal{M}$, $E_k(D_k(m)) = E(D(m, k), k) = m = D(E(m, k), k) = D_k(E_k(m))$. Therefore E_k, D_k are functions which are each other's inverses. Therefore they are bijective. \square

This also makes it clear that for each key, each plaintext is associated with one ciphertext. If we did not have this then the decryption of a ciphertext could result in multiple plaintexts, making the cipher effectively useless. Note that only holds for a fixed key: a plaintext can be encrypted to a number of different ciphertexts if you vary the key. If we keep this key secret, then this one to many association between plaintexts and ciphertexts under different keys leads to one of the important properties we want from a block cipher: confidentiality.

3.2 Confidentiality, OTP and Xor

Block Ciphers can provide many services to their users, but the one we are going to examine in this project is *Confidentiality*: ensuring that if an adversary can intercept ciphertexts then that adversary cannot reasonably recover any information about the corresponding plaintexts. We use the term *adversary* to represent an individual who wants to uncover the contents of encrypted messages. This is used to model potential attacks on our current cryptographic scheme.

This idea of confidentiality was first strongly formalised by Shannon with his concept of *Perfect Secrecy* [6].

Definition 22 (Shannon's Perfect Secrecy). *Let us fix m_1, m_2, c as arbitrary elements of \mathcal{M} . A cipher is said to have **perfectly secrecy** if it has either of these equivalent properties holds for k chosen uniformly over the key space:*

1. $Pr[E(m_1, k) = c] = Pr[E(m_2, k) = c]$
2. $Pr[E(m, k) = c] = \frac{1}{|\mathcal{M}|}$

Essentially, if an adversary is given a ciphertext encrypted with a cipher with perfect secrecy, that adversary can recover no information about the plaintext. This is because this plaintext is indistinguishable from a plaintext randomly picked from \mathcal{M} . So even if an adversary did manage to recover a ciphertext, without knowing the secret key the confidentiality of the message is retained.

An example of a cipher with perfect secrecy is the *One Time Pad (OTP)*. This can be seen as a block cipher with one round. Before we can define OTP, we need to define what is meant by a *binary xor*.

Definition 23 (Binary Xor). *The binary exclusive or (otherwise known as **xor**) is the function*

$$\oplus : \mathbb{Z} \times \mathbb{Z} \rightarrow \mathbb{Z}$$

If we take $a, b \in \mathbb{Z}$ with bit-strings a', b' of length n and define x_i for $0 \leq i < n$ as the i th digit of the n bit string x , then we can define $a \oplus b = c$ where

$$c_i = a'_i + b'_i \text{ mod } 2$$

Example 24. *Here is the result of \oplus on all the single bit numbers:*

$$1 \oplus 1 = 0 = 0 \oplus 0, \quad 1 \oplus 0 = 0 = 0 \oplus 1 \tag{1}$$

If we consider 1 to be the value representing True and 0 to be the value representing False, then \oplus corresponds to the exclusive or function (hence the name, xor). If we take boolean values a, b we have

$$a \oplus b = (\neg(a \wedge b)) \wedge (a \vee b)$$

It is worth noting some useful properties of xor as we will be using in the following sections.

Lemma 25 (Xor Properties). \oplus is commutative, associative, each element is its own inverse and the identity of \oplus is 0.

Proof. Since \oplus operates on each bit of a bit string independently, it suffices to show that all of these properties hold for one bit numbers. If one takes $a, b, c \in \mathbb{Z}_2$ it is easy to show that all of these hold from the calculations in Equation 1. \square

With this defined, we can easily define OTP through its encryption and decryption algorithms.

Definition 26. For $m, k \in \mathcal{M} = \mathcal{K}$ being all bit strings of a certain maximum bit length b , we define the **One Time Pad** as the block cipher with encryption and decryption functions

$$E_{OTP}(m, k) = m \oplus k, \quad D_{OTP}(c, k) = c \oplus k$$

Since $(m \oplus k) \oplus k = m \oplus (k \oplus k) = m$, it is easy to verify that for a fixed key, the encryption under OTP through E_{OTP} is reversed by D_{OTP} .

Theorem 27. OTP has perfect secrecy.

Proof. Take $m, c \in \mathcal{M}$. The probability that $E(m, k) = c$ over all choices of k from \mathcal{M} will be equal to the number of keys that give $E(m, k) = c$ divided by the total number of keys. If $E(m, k) = c$ then

$$m \oplus k = c \Rightarrow m \oplus c = k$$

Therefore, there is only one possible key under which the encryption of m will result in c . Therefore the probability that $E(m, k) = c$ is for some random k is $\frac{1}{|\mathcal{K}|}$. Therefore, the probability is equal for any choice of m , i.e. for $m' \in \mathcal{M}$ the probability that $E(m', k) = c$ is the same as for $E(m, k) = c$. This means that OTP has perfect secrecy. \square

We use a process very similar to OTP for key addition in the cipher we will be analysing - Unsub A. In the next section, we discuss circumstances where perfect secrecy is not enough.

3.3 Chosen-plaintext attack and Kerckhoffs's principle

A block cipher needs to maintain the confidentiality of many ciphertexts all encoded with the same key. While OTP has perfect secrecy, this does not mean OTP maintains confidentiality no matter how it is used or over all types of attacks.

Example 28. *If we were to encode two messages with the same key, i.e. take $k \in \mathcal{K}, m_0, m_1 \in \mathcal{M}$ with*

$$E_{OTP}(m_0, k) = c_0, \quad E_{OTP}(m_1, k) = c_1$$

then if an adversary recovers c_0 and c_1 then they can calculate $c_0 \oplus c_1 = (m_0 \oplus k) \oplus (m_1 \oplus k) = m_0 \oplus m_1$. While they have recovered no information about the key, they have definitely recovered some information about the plaintexts so perfect secrecy is broken.

If we give the adversary more power in that we give them the ability not only to look at some ciphertexts but give them ability to choose some plaintexts and get their corresponding ciphertexts encrypted with the secret key, then OTP used with the same key twice has no confidentiality.

Example 29. *If an adversary requests the encryption of m from \mathcal{M} and receives $E_{OTP}(m, k) = c$ with $k \in \mathcal{K}$ not being revealed, then the adversary can simply calculate $c \oplus m = (m \oplus k) \oplus m = k$ from the information given.*

In Example 28, in order to recover information about the plaintexts we took the xor of them both and examined that value. If we view the xor of the two values as a difference, then we can view this as a proto-differential attack on OTP: we used the differences between input values in order to determine secret information. This demonstrates the power of taking differences with \oplus which we will expand upon later with much more complex block ciphers.

If we allow the attacker to choose plaintext values and get the corresponding ciphertexts before undertaking an attack, then the attacker is undertaking what is known as a *Chosen-plaintext attack*.

Definition 30. A **chosen-plaintext attack** (C-PTA) occurs in two steps.

1. An adversary can request the encryption of n plaintexts m_0, \dots, m_n with $m_i \in \mathcal{M}$ and receive the corresponding ciphertexts c_0, \dots, c_n under the secret key k , with $E(m_i, k) = c_i$.
2. After this, the adversary is allowed to choose a plaintext m , with $m \neq m_i$ for $i \in \{0, 1, \dots, n-1, n\}$. They receive two ciphertexts c, c' . Of these ciphertexts, one is the encryption of m under the secret key k and the other is a random value chosen from \mathcal{M} .

If the adversary can differentiate between whether c or c' is $E(m, k)$ with a technique that requires fewer than $|\mathcal{K}|$ guesses for k then say that a cipher is insecure under C-PTA.

There is an alternate definition for this that involves non-negible probability but that is far harder to work with than our approach above.

After fully defining C-PTA, we have what we call the parameters of our attack. This defines the abilities we have (obtaining those plaintext-ciphertext pairs) and the criteria for success (differentiating between a random and a non-random ciphertext with fewer guesses than a brute force attack).

In our analysis we will be assuming that the adversary has the powers of a chosen plaintext attack, i.e. the adversary can get a certain number of plaintext-ciphertext pairs from the cipher. We will also be assuming *Kerckhoff's Principle*: a cipher should be secure if an adversary knows every detail about it besides its key. As most of the very prominent block ciphers such as AES and DES have all their implementation details in the public domain, this is not an unreasonable principle to hold. This means that the adversary can choose certain plaintexts based on the properties of the cipher. All of the specifications for our cipher Unsub A are explained in full in Appendix A in accordance with this principle.

In the following sections, we are not going to explicitly refer to C-PTA to determine whether our attacks are successful or not. This is because successful differential attacks should recover either the whole or partial key required to encrypt or decrypt values more efficiently than a brute force attack. They can then recover the rest of the key if needed through a brute force attack and simply decrypt both ciphertexts given to it, showing the cipher to be insecure under C-PTA.

4 Basic Elements of Differential Cryptanalysis

In this section, we are going to start the work of showing that UnsubA (as described in Appendix A) is insecure under C-PTA using differential cryptanalysis. The traditional approach to teach this is to take ciphers of increasing complexity and build up the techniques of differential cryptanalysis by breaking each in turn. We are going to take a slightly different approach here: instead of looking at progressively more complex ciphers, we are going to begin by looking at UnsubA and determine how much information we can deduce about the difference between values of the partially encrypted plaintext after various rounds if we know the difference between the initial plaintext values.

4.1 Outline of an Attack using Differential Cryptanalysis

As mentioned earlier, differential cryptanalysis studies differences. In its simplest form, it examines how differences in plaintext pairs affect the differences in the corresponding ciphertexts. Below we formally define what we mean by a *difference*.

Definition 31. For $x, y \in \mathcal{M}$ we have that the **difference** of x and y with respect to a function \otimes is

$$\Delta(x, y) = x \otimes y^{-1}$$

We choose this difference function according to what will make the cryptanalysis easiest. This means that we usually define \otimes to be the operation used for key addition as this means that if we take the differences of x, y after key addition we have that if the key addition algorithm is associative and commutative

$$\Delta(x \otimes k, y \otimes k) = (x \otimes k) \otimes (y \otimes k)^{-1} = x \otimes y^{-1} = \Delta(x, y)$$

For Unsub A, this is binary xor which we denote by \oplus . Since this project is about primarily analysing Unsub A, we will be developing our analysis with \oplus rather than the general \otimes . Please note that with this definition of difference above one can reword our results to be more general with little difficulty. We do this because for \oplus , $y^{-1} = y$ for all $y \in \mathcal{M}$, simplifying our calculations somewhat.

In order to motivate why we are trying to find these differences, we are going to assume we have already found these differences and show how this means we can now find the key for the cipher. We will give the attack a comprehensive treatment in Section 6, this section is just to give some intuition so that we do not simply leap in finding differences blindly. We are also going to focus on a generic attack which attempts to find the entire last round key - later we will go into detail on how to find just certain bits of the key.

In our SPN cipher UnsubA, we have 5 rounds of encryption all consisting of a key addition, substitution and permutation. After these 5 rounds, there is a final key addition

done. Let us assume by analysing a r -round iterated block cipher we have found that if we take plaintext pairs $m = (m_0, m_1)$ with differences of $\Delta m = \alpha_0$ and encrypt them with all but the last substitution-permutation and key addition round of our cipher that we get that the corresponding output pairs will have difference $\Delta l = \alpha_{r-1}$ with probability p . Since this is a C-PTA, let us take N plaintext pairs with difference α_0 giving us N many ciphertexts.

Now, for every ciphertext pair (c_0, c_1) corresponding to one of our plaintext pairs (m_0, m_1) with $\Delta(m_0, m_1) = \alpha_0$ and for every k that could possibly be the last key used in our cipher, we are going to partially decrypt c_0, c_1 using our k to get l_0, l_1 , the potential values of m_0, m_1 after $r - 1$ rounds of encryption. While we do not know what l_0 or l_1 is, we do know that $\Delta(l_0, l_1) = \alpha_{r-1}$ with probability p when partially decrypted with the correct value for our last round key. Therefore, if $\Delta(l_0, l_1)$ does equal α_{r-1} then k is a possible key value, so we record it.

After we have done this for every key and ciphertext pair, we look at the frequency that each key k was recorded as a possible key for our last round. We can then guess the value of this round key based on the expected frequency that it would occur. This calculation is done fully in Section 6 but we expect the key that was suggested with frequency pN over the N pairs to be the correct key for large enough p, N .

As you can see, we have used our ability to predict the likelihood of the difference of an output pair given the difference of a plaintext pair to get us part of the key of a block cipher. In the next section, we will be looking at the rounds of a block cipher in order to build up a framework for us to be able to start making these predictions.

4.2 Differential Cryptanalysis on One Round

Before being able to predict differences across $r - 1$ rounds, we first need to be able to predict them over a single round. Here we analyse how we can do this for an SPN block cipher like Unsub A.

4.2.1 Key Addition

Since we are dealing with a C-PTA, let us pick two plaintexts m_0, m_1 from \mathcal{M} and let us assume there is a key $k \in \mathcal{K}$ that is not known to us that determines K^i , a function representing the key addition step in round i . Let us set $v_j = K^i(m_j), w_j = S(v_j), u_j = P(w_j)$ for $j = 0, 1$ and S, P being the substitution and permutation steps respectively for our iterated block cipher. Now, knowing m_0, m_1 , what can we say about $\Delta(u_1, u_2)$, the difference after a round?

As shown in Theorem 27, if m_j and k are picked with uniform random probability then v_j is a uniform random variable on \mathcal{M} . Therefore without knowing k , we cannot have an any guess better than random for the value of v_j . However, we know for UnsubA that

the key addition is done with xors which means that although we cannot accurately guess what v_j is, we can calculate the value of $v_0 \oplus v_1$:

$$v_0 \oplus v_1 = (m_0 \oplus k) \oplus (m_1 \oplus k) = m_0 \oplus m_1$$

Note that this means that we can effectively ignore the key addition stage if we are considering differences as we can predict the difference between two output values given the difference between the inputs.

Example 32. Suppose we choose the ciphertexts $m = (m, m') = (0xa, 0x5)$. These have a difference of $\Delta m = 0xf$ which means by the above, the difference of the corresponding ciphertext will also be $0xf$. To test this, let us take the key to be $k = 0x1$. Then

$$m_1 \oplus k = 0xa \oplus 0x1 = 0xd, m_2 \oplus k = 0x5 \oplus 0x1 = 0x5$$

This means that the corresponding ciphertexts $c = (c, c') = (0xd, 0x5)$ which one can verify have a difference of $\Delta c = 0xf$.

Now, let us suppose we did not have our key, we only had m and c . Could we recover some key bits? Well, we could just calculate:

$$m_1 \oplus k = c_1 \Rightarrow m_1 \oplus c_1 = 0xa \oplus 0xd = 0x1 = k$$

4.2.2 Substitution

Now, with $v_0 \oplus v_1$ known we want to find $w_0 \oplus w_1 = S(v_0) \oplus S(v_1)$. While we have $v_0 \oplus v_1$ and all the values of $S(m)$ for $m \in \mathcal{M}$, we do not have v_j so we cannot calculate $\Delta S(v_j)$. However, we can construct a *difference distribution table* to give us a value to expect.

Definition 33 (Difference Distribution Table). A **difference distribution table** D is defined as a table where the cell in the a th row and the b th column $D_{a,b}$ represents the frequency that $S(p_0) \oplus S(p_1) = b$ for all $p_0, p_1 \in \mathcal{M}$ such that $p_0 \oplus p_1 = a$.

To calculate a row a for a difference distribution table, one first initialises all values in the row to zero. Then one iterates through all possible input values for the sbx setting them to p_0 , getting p_1 by noting that

$$p_0 \oplus p_1 = a \Rightarrow p_1 = a \oplus p_0$$

With this we get every possible pair with difference a . We then simply calculate the difference of their substituted values $S(p_0) \oplus S(p_1) = b$ and increment the cell $D_{a,b}$ which this difference corresponds to.

Using D and $v_0 \oplus v_1$, we can calculate the possible values of $w_0 \oplus w_1$ by looking at which cells in the $v_0 \oplus v_1$ row of D are non-zero. We can also use D to calculate the likelihood of values of differences after going through the substitution step by dividing each value in D by $|\mathcal{M}|$. This is proved below.

Lemma 34. *Given a pair with the difference a the probability that the corresponding pair of outputs after going through an sbox S will have a difference of b is*

$$\frac{D_{a,b}}{|\mathcal{M}|}$$

Proof. Lets take $\Omega = \mathcal{M}$ and assume we have a uniform probability distribution. Let us define the random variable X on \mathcal{M} such that for every $x \in \mathcal{M}$,

$$X(x) = S(x) \oplus S(x \oplus a)$$

This of course is the output difference of a pair of outputs with difference a of which x is a member. With this, the probability that for an input difference a we will get an output difference of b is

$$Pr[X = b] = Pr[A] \text{ with } A := \{y \in \mathcal{M} | S(y) \oplus S(y \oplus a) = b\}$$

The cardinality of A will be the number of elements of \mathcal{M} which when in a pair of difference a gives a difference of b . But this is exactly $D_{a,b}$. Therefore, as we are dealing with a uniform probability distribution

$$Pr[X = b] = \frac{D_{a,b}}{|\mathcal{M}|}$$

as desired. □

Using our difference table D , we can now calculate the most likely value for $w_0 \oplus w_1$ which holds with some probability p . What can we say about this probability? We can say that if $m_0 = m_1$ this probability will be 1 as, of course, $S(m_0) \oplus S(m_0) = 0$ will always hold. This means that the first row and column of the table will be 0 except for the the cell in the first row and column which will be $|\mathcal{M}|$. We can also note that the rows each add up to $|\mathcal{M}|$ as for each $m \in \mathcal{M}$ and difference d there is exactly one $m' \in \mathcal{M}$ such that $\Delta(m, m') = d$. This means that each input difference will have $|\mathcal{M}|$ output differences giving us that number in each row. Putting these together means that for every row other than the row for 0 input differences, there must be at least one cell with a value larger than or equal to 2 which means we can always find a value which is more likely than $\frac{1}{|\mathcal{M}|}$, i.e. picking an output difference at random.

Note that normally the substitution layer of a block cipher is often divided into a number of identical sboxes as done in UnsubA. If this is the case, one usually calculates one difference distribution table across one sbox and then uses that to calculate the probability a difference holds across the entire substitution layer by noting that every sbox's will act its input values independently of the other sboxes. We look at this in the example below.

Example 35. Let us take $\Delta m = 0x222120 = 0b1000100010000100100000$. To look this up in an sbox, we have to split $0x222120$ into the four 8-bit binary numbers $(m_0, m_1, m_2, m_3) = (0b0, 0b100002, 0b100001, 0b100000)$, each of which go to a different s-box. To find the output difference that these input differences m_i will most likely lead to, we find the d_i in our difference distribution table with the largest value, that is to say

$$d_i = \max\{D_{m_i, d}\} \text{ for } d = 0 \text{ to } d = 2^8$$

If we do this, we get the following for each sbox input m_i above:

Input Difference	Most likely output difference	Probability output difference is achieved
0x0	0x0	0x1
0x22	0x2	≈ 0.0312500
0x21	0x62	≈ 0.0312500
0x20	0x2	≈ 0.0859375

Since all these inputs are independent, the most likely output difference for the input of $0x222120$ is $0x26202$ with probability approximately $0.03125 \times 0.03125 \times 0.0859375 \times 1 = 8.392333984375 \times 10^{-5}$.

4.2.3 Permutation

Now that we have a value for $w_0 \oplus w_1$ which holds with some probability p , we want to calculate $u_0 \oplus u_1$ with as much certainty as possible. Now, since in SPN block ciphers the permutation layers are linear with respect to \oplus , we can simply calculate this directly from $w_0 \oplus w_1$

$$u_0 \oplus u_1 = P(w_0) \oplus P(w_1) = P(w_0 \oplus w_1)$$

which will hold again with probability p .

If the permutation is not linear, then one can simply apply the techniques we developed in the section on non-linear substitution boxes to them in order to get a probabilistic value for $u_0 \oplus u_1$.

For Unsub A, it is quite easy to prove this is linear over \oplus .

Lemma 36. *Unsub A's permutation layer is linear with respect to binary xor*

Proof. Take $x, y \in \mathcal{M}$ and let us define σ as a permutation on the set of 5 bit numbers as defined in Table 2a which defines the permutation layer of Unsub A. Then the i th bit of the permutation of x is $x_{\sigma(i)}$ where x_j is the j th bit of x . With this in mind, let us consider the i th bit of $P(x) \oplus P(y)$. This will be

$$\begin{aligned}
(P(x) \oplus P(y))_i &= x_{\sigma(i)} \oplus y_{\sigma(i)} \\
&= (x_i \oplus y_i)_{\sigma(i)} \\
&= P(x \oplus y)_i
\end{aligned}$$

This holds as \oplus and P are bit-wise operations so we can consider their operations on each bit separately. This means that P is linear with respect to \oplus . \square

4.2.4 Characteristics

Without expressive notation, it becomes quite difficult to follow the logic behind the differential analysis when we get to multiple rounds. In order to aid clarity, we can simplify the notation we used in the previous sections by defining *characteristics*.

Definition 37. *The pair $(\Delta\alpha_0, \Delta\alpha_1) \in \mathcal{M} \times \mathcal{M}$, with $\Delta\alpha_0$ being an input difference and $\Delta\alpha_1$ being an expected output difference after some operation O , is known as a **differential characteristic** across O . This characteristic holds with a probability p .*

This will be denoted by $\Delta\alpha_0 \xrightarrow{O} \Delta\alpha_1$

We can now represent the result of our calculations using this new notation. Let us first set the value of our calculations to variables:

$$m = (m_0, m_1), v = (v_0, v_1), w = (w_0, w_1), u = (u_0, u_1)$$

Then we can represent our calculations we did in the previous sections by saying

- $\Delta m \xrightarrow{K^i} \Delta v$ with probability 1
- $\Delta v \xrightarrow{S} \Delta w$ with probability p
- $\Delta w \xrightarrow{P} \Delta u$ with probability p

To further simplify our notation, we can define an operation \mathcal{R}^i to represent the operations of our cipher for round i with

$$\mathcal{R}^i = P(S(K^i(a)))$$

We will often leave out the round number i and allow the user to infer the round from the context we use it in.

This means that our entire calculation we performed in this section amounts to $\Delta m \xrightarrow{\mathcal{R}^i} \Delta u$ which holds with probability p .

Example 38. Let us look at a characteristic starting with an input difference of 0x32 and work out step by step what this will give us over a round. With key addition, we have that $0x32 \xrightarrow{K^1} 0x32$ with probability 1. To get the expected value over the sboxes, we need to split up 0x32 into 8-bit pieces and look up these pieces in a difference distribution table giving us the below:

Input Difference	Most likely output difference	Probability output difference is achieved
0x0	0x0	0x1
0x0	0x0	0x1
0x0	0x0	0x1
0x20	0x2	≈ 0.0859375

This means that $0x32 \xrightarrow{K^i} 0x32 \xrightarrow{S} 0x2$ holds with probability ≈ 0.0859375 . Now, as we saw, since P is linear with respect to \oplus , we have that $0x32 \xrightarrow{K^i} 0x32 \xrightarrow{S} 0x2 \xrightarrow{P} 0x2$ which is equivalent to saying $0x32 \xrightarrow{\mathcal{R}^i} 0x2$ which holds with probability ≈ 0.0859375 . It turns out that for UnsubA, this is one of the single round characteristics that holds with the largest probability - we will be examining it in more depth later.

Using the above, we can now determine with some probability what the difference of a input pair will be after a round of a block cipher. What we are going to do is extend this approach by joining the characteristics of rounds together, allowing us to predict the difference of the ciphertexts when given a pair of plaintexts with a certain difference.

4.3 Multiple Rounds

We know how to create characteristics for one round, but in order for this to be useful for even our toy cipher we need some way to predict the output difference after multiple rounds. To do this, we extend the concept of a characteristic to cover multiple rounds. This is called an *s-round characteristic*.

4.3.1 S-round Characteristics

Definition 39. If we denote the output difference of a round as Δc_i with the input difference denoted as Δm then an **s-round characteristic** is a list of differences represented as the tuple $(\alpha_0, \dots, \alpha_s)$ where α_i is the anticipated value of Δc_i with $\Delta c_0 = \Delta m$.

Given an s-round characteristic of this form, what we want to calculate is the probability that the output differences for Rounds 1 through i of the encryption are what we predict. Formally, what we want is

$$Pr[\Delta c_i = \alpha_i, \Delta c_{i-1} = \alpha_{i-1}, \dots, \Delta c_1 = \alpha_1 | \Delta m = \alpha_0]$$

with this probability being taken over every possible choice of key from \mathcal{K} and plaintext pair with that difference from \mathcal{M} ³. To calculate this, we are going to develop the idea of a *Markov Cipher*.

4.3.2 Markov Chains and Ciphers

Since we have developed a decent way of calculating the probabilities for single round characteristics, we would ideally want to be able to calculate the probability of an s-round characteristic using single round characteristic. We want to be able to “chain” together single round characteristics, calculating the probability for each successive characteristic by simply considering the preceding one.

It turns out what we want our chain of s-round characteristics to be is a *homogeneous Markov chain*.

Definition 40. A sequence of random variables c_0, c_1, \dots, c_r is called a **Markov chain** if the probability of each c_i is only dependent on the previous value c_{i-1} . That is, for $0 \leq i < r$,

$$Pr[c_{i+1} = \alpha_{i+1} | c_i = \alpha_i, c_{i-1} = \alpha_{i-1}, \dots, c_1 = \alpha_1, c_0 = \alpha_0] = Pr[c_{i+1} = \alpha_{i+1} | c_i = \alpha_i]$$

If we have that the value of $Pr[c_{i+1} = \alpha_{i+1} | c_i = \alpha_i]$ is not dependent on the choice of i for all α_j then we have a **homogeneous Markov chain**.

Let us assume that we can prove that $\Delta c_0, \dots, \Delta c_r$ forms a homogeneous Markov chain. What we want to show is that this gives us the ability to calculate the probability of our s-round characteristic by simply multiplying the probability of 1-round characteristics together.

Theorem 41. Given $\Delta m = \Delta c_0, \dots, \Delta c_r$ is a homogeneous Markov chain, we have that the probability that the r-round characteristic $(\alpha_0, \alpha_1, \dots, \alpha_{r-1}, \alpha_r)$ holds with probability

$$Pr[\Delta c_r = \alpha_r, \Delta c_{r-1} = \alpha_{r-1}, \dots, \Delta c_1 = \alpha_1 | \Delta m = \alpha_0] = \prod_{i=1}^r Pr[\Delta c_i = \alpha_i | \Delta m = \alpha_{i-1}]$$

Proof. Let us do this by induction on r , the number of rounds in the characteristic. For $r = 2$, this is obvious. Let us assume this is true for k rounds. Then we want to show that

$$Pr[\Delta c_{k+1} = \alpha_{k+1}, \Delta c_k = \alpha_k, \dots, \Delta c_1 = \alpha_1 | \Delta m = \alpha_0] = \prod_{i=1}^{k+1} Pr[\Delta c_i = \alpha_i | \Delta m = \alpha_{i-1}]$$

³Note that this does not model the real life case exactly where we have a single key which is fixed and unknown. We address this later by considering the Principle of Stochastic Equivalence.

To do this, first let us expand the Markov Chain Definition. Since $\Delta c_0, \dots, \Delta c_{k+1}$ is a Markov Chain, we have that

$$\begin{aligned} & Pr[\Delta c_{k+1} = \alpha_{k+1} | \Delta c_k = \alpha_k, \Delta c_{k-1} = \alpha_{k-1}, \dots, \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] \\ &= Pr[\Delta c_{k+1} = \alpha_{k+1} | \Delta c_k = \alpha_k] \end{aligned}$$

If we expand these out using the definition of conditional probability, we get that

$$\begin{aligned} & \frac{Pr[\bigcap_{i=0}^{k+1} \Delta c_{i+1} = \alpha_{i+1}]}{Pr[\bigcap_{j=0}^k \Delta c_j = \alpha_j]} \\ &= \frac{Pr[\Delta c_{k+1} = \alpha_{k+1} \bigcap \Delta c_k = \alpha_k]}{Pr[\Delta c_k = \alpha_k]} \end{aligned}$$

Then, if we multiply both sides of the equality by

$$\frac{Pr[\bigcap_{j=0}^k \Delta c_j = \alpha_j]}{Pr[\bigcap_{j=0}^k \Delta c_j = \alpha_j]}$$

we get that

$$Pr[\bigcap_{i=0}^{k+1} \Delta c_{i+1} = \alpha_{i+1}] = \frac{(Pr[\bigcap_{j=0}^k \Delta c_j = \alpha_j])(Pr[\Delta c_{k+1} = \alpha_{k+1} \bigcap \Delta c_k = \alpha_k])}{Pr[\Delta c_k = \alpha_k]}$$

If we expand out the expression for the probability that our $k + 1$ round characteristic holds, we get

$$\begin{aligned} & Pr[\Delta c_{k+1} = \alpha_{k+1}, \Delta c_k = \alpha_k, \dots, \Delta c_1 = \alpha_1 | \Delta c_0 = \alpha_0] \\ &= \frac{Pr[\bigcap_{i=0}^k \Delta c_{i+1} = \alpha_{i+1}]}{Pr[\Delta c_0 = \alpha_0]} \\ &= \frac{Pr[\bigcap_{j=0}^k \Delta c_j = \alpha_j] Pr[\Delta c_{k+1} = \alpha_{k+1} \bigcap \Delta c_k = \alpha_k]}{Pr[\Delta c_k = \alpha_k] Pr[\Delta c_0 = \alpha_0]} \\ &= Pr[\Delta c_{k+1} = \alpha_{k+1} | \Delta c_k = \alpha_k] \prod_{i=0}^k Pr[\Delta c_1 = \alpha_i | \Delta c_0 = \alpha_{i-1}] \text{ (Inductive hypothesis)} \\ &= \prod_{i=0}^{k+1} Pr[\Delta c_1 = \alpha_i | \Delta c_0 = \alpha_{i-1}] \text{ (Property of homogenous Markov Chains)} \end{aligned}$$

Since $\Delta m = \Delta c_0$, we have that by mathematical induction, this result holds. \square

This means that if our chain of characteristics for our cipher is a homogenous markov chain, then by analysing the difference distribution table on one sbox and deriving the probabilities for one-round characteristics we have enough to calculate the probability of any s-round characteristic if we can show that the characteristics form a homogeneous Markov chain.

The next natural question to ask is what types of ciphers produce characteristics which are homogeneous Markov chains. We would expect such ciphers to necessarily have the quality that the value of the pair does not influence the probability taken over all the keys that a one-round characteristic holds for a given difference. If it did, then any individual characteristic Δc_i would be dependent on more than just Δc_{i-1} which means that the characteristic could not form a Markov chain. Ciphers that have this quality are called *Markov Ciphers*. More formally:

Definition 42. *An iterated cipher is called a **Markov Cipher** if $Pr[\Delta c_1 = \alpha_1 | \Delta c_0 = \alpha_0, c_0 = \gamma]$ is independent of γ for all α_i when the round key k is a uniform random variable.*

It turns out that this quality by itself, with some stipulation on the keys, is enough to ensure that characteristics are homogeneous Markov chains as the theorem below shows.

Theorem 43. *If we have a sequence of differences $\Delta m_0 = \Delta c_0, \Delta c_1, \dots, \Delta c_r$ created by an r -round Markov Cipher with the r round keys generated independently and uniformly at random, then that sequence of differences is a homogeneous Markov chain.*

Proof. This proof is courtesy of Lai and Massey[7].

Since we are looking at iterated ciphers, to show that $\Delta c_0, \Delta c_1, \dots, \Delta c_r$ is a Markov Chain, we only need show that

$$Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] = Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1]$$

To prove our result, we note that

$$\begin{aligned} & Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] \\ &= \sum_{\gamma} Pr[\Delta c_2 = \alpha_2, c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] \\ &= \sum_{\gamma} Pr[c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, c_1 = \gamma, \Delta c_0 = \alpha_0] \end{aligned}$$

Now, if we specify Δc_1 as well as c_1 then that determines Δc_2 completely as this is an iterated cipher - the input from the end of the last round is given to the next, so the result

of the $(n - 1)$ th round determines the n th one. This means that

$$\begin{aligned} & \sum_{\gamma} Pr[c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, c_1 = \gamma, \Delta c_0 = \alpha_0] \\ &= \sum_{\gamma} Pr[c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, c_1 = \gamma] \end{aligned}$$

Now, since we are dealing with a Markov Cipher, we have that

$$Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, c_1 = \gamma] = Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1]$$

This gives us that

$$\begin{aligned} & \sum_{\gamma} Pr[c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1, c_1 = \gamma] \\ &= \sum_{\gamma} Pr[c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0] Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1] \\ &= Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1] \underbrace{\sum_{\gamma} Pr[c_1 = \gamma | \Delta c_1 = \alpha_1, \Delta c_0 = \alpha_0]}_{=1} \\ &= Pr[\Delta c_2 = \alpha_2 | \Delta c_1 = \alpha_1] \end{aligned}$$

Since we are dealing with an iterated cipher, we immediately obtain that this Markov Chain is homogenous. This is the case as the round function does not vary except for round key additions which do not affect the differences of pairs by our choice of the difference operator. Therefore since we proved it holds over one round, it can hold over any of the rounds. Thus the result is proven. \square

For our analysis of Unsub A, we would like to be able to use the tools we have developed above to calculate the probabilities of s -round characteristics. In Section 5 we will be discussing how to find such characteristics using the probabilities of one round characteristics, but before we can do that we have to show that UnsubA is a Markov Cipher.

Theorem 44. *Unsub A is a Markov Cipher.*

Proof. Let us fix some $\alpha_0, \alpha_1 \in \mathcal{M}$. Let us take $\gamma = (m_1, m_2)$ where m_1, m_2 are arbitrary elements of \mathcal{M} st. $\Delta c_0 = \alpha_0$. As we have discussed above, input differences are not preserved through non-linear substitution layers like in UnsubA. Therefore $\Delta c_1 = \alpha_1$ is not independent of $\Delta c_0 = \alpha_0$. To see if Δc_1 is independent of $c_0 = \gamma$, we need to show that the values of γ have no effect on their corresponding ciphertexts after one round of encryption.

Let us choose the round key k uniformly from \mathcal{K} . Then we have that the key addition step is OTP encryption:

$$u_i = K^1(m_i) = m_i \oplus k, i = 1, 2$$

This means, as OTP is perfectly secure (see Theorem 27), we have that u_1, u_2 are equal to any element in \mathcal{M} with probability $\frac{1}{|\mathcal{M}|}$. This means that they might as well have been picked randomly from \mathcal{M} . Since the substitution and permutations layers are permutations, the probability that $P(S(K^1(m_i)))$ for $i = 1, 2$ is equal to any element in \mathcal{M} is still $\frac{1}{|\mathcal{M}|}$. Therefore the chance that

$$P(S(K^1(m_1))) \oplus P(S(K^1(m_2))) = m'$$

for arbitrary $m' \in \mathcal{M}$ is $\frac{1}{|\mathcal{M}|}$. In other words, the initial choice of γ has no impact on the value of Δc_1 . Therefore

$$Pr[\Delta c_1 = \alpha_1 | \Delta c_0 = \alpha_0, c_0 = \gamma] = Pr[\Delta c_1 = \alpha_1 | \Delta c_0 = \alpha_0]$$

□

4.3.3 Differentials

In the previous section, we fully specified the value expected at the end of each round of our cipher using s-round characteristics such as

$$\alpha_0 \xrightarrow{\mathcal{R}} \alpha_1 \xrightarrow{\mathcal{R}} \dots \xrightarrow{\mathcal{R}} \alpha_{s-1} \xrightarrow{\mathcal{R}} \alpha_s$$

We call pairs with input difference α_0 right pairs with respect to a characteristic if they follow the characteristic and wrong pairs if they do not. For example, if we take the sequence of values $(\alpha_0, \gamma, \alpha_2, \alpha_3, \dots, \alpha_{s-1}, \alpha_s)$ this is considered a wrong pair if $\gamma \neq \alpha_1$.

However, for the attack, we do not really care what the intermediary values are in the cipher. To recover key bits, all we need to know is the probability that we get an output difference Δl given an input difference Δm . If our input difference was α_0 and our end predicted difference was α_s then we do not just want to know the probability that the above characteristic holds but the probability for all characteristics of the form

$$\alpha_0 \xrightarrow{\mathcal{R} ?} \dots \xrightarrow{\mathcal{R} ?} \alpha_s$$

holds where $?$ denotes an intermediary value in the cipher which can be any value that leads to the desired α_s end difference. This means that we cannot just consider right pairs but any wrong pair with the same starting and ending value as our characteristic. This is captured by the concept of a *differential*.

Definition 45. A *s-round differential* is a pair of differences (α_0, α_s) where α_s is the expected value of Δc_s given the input difference $\Delta m = \alpha_0$.

Just as with s-round characteristics, we can calculate differentials of Markov Ciphers with one round characteristics. But before we do that, we have to tackle the issues of key equivalence.

We say that a a plaintext pair is a *right pair* with respect to a s -round differential and a fixed key if when it is encrypted using the key, the difference predicted by the differential is the difference of its values after s rounds.

An r -round characteristic can only be a homogeneous Markov chain if the choice of key does not affect the value of the difference. While this is true in some ciphers, it is definitely not guaranteed. When performing our analysis, we generally assumed that the key is randomly and uniformly picked from the key space. However, when actually undertaking the attack, one will be encrypting plaintext pairs with a fixed key. There is nothing to say that the probability that the difference holds over all keys chosen randomly will be equal to the probability that difference holds over some fixed key [7].

In order to allow us to continue our analysis, we assume that keys will all behave roughly in the same manner in a hypothesis called the *hypothesis of stochastic equivalence*.

Definition 46. The *hypothesis of stochastic equivalence* says that for almost all high probability r -round differentials (α, β) on a cipher encrypting with key k , we have

$$Pr[\Delta c_r = \beta | \Delta m = \alpha, k = k'] \approx Pr[\Delta c_r = \beta | \Delta m = \alpha]$$

holds for a large number of key values $k' \in \mathcal{K}$.

We will implicitly assume this is the case as for Unsub A for the rest of the project in order to move forward in our analysis.

With this in mind, we can calculate the value of a differential.

Theorem 47. The probability a differential (α_0, α_s) holds in a Markov cipher is the sum of the probabilities of all the s -round characteristics starting with α_0 and ending with α_s .

To put this formally, if the largest value a difference can take is $\max(\mathcal{M})$ then

$$Pr[\Delta c_s = \alpha_s | \Delta m = \alpha_0] = \sum_{j=1}^{s-1} \sum_{\alpha_j=0}^{\max(\mathcal{M})} \prod_{i=1}^s Pr[\Delta c_1 = \alpha_i | \Delta m = \alpha_{i-1}]$$

Proof. We will show this through induction. For a $s = 1$ round differential, we are essentially looking at the probability that the characteristic $\alpha_0 \xrightarrow{\mathcal{R}} \alpha_1$ holds which is $Pr[\Delta c_1 = \alpha_1 | \Delta m = \alpha_0]$. Therefore the base case is true.

Let us assume that this formula holds for $s = k$ round differentials. Then let us consider the $k + 1$ round differential (β_0, β_{k+1}) . Let us take a $\gamma \in \mathcal{K}$ and examine the probability for the k round differential (β_0, β_k) . By our inductive hypothesis the probability this holds is

$$Pr[\Delta c_k = \beta_k | \Delta m = \beta_0] = \sum_{j=1}^{k-1} \sum_{\beta_j=0}^{\max(\mathcal{M})} \prod_{i=1}^k Pr[\Delta c_1 = \beta_i | \Delta m = \beta_{i-1}]$$

This means that

$$Pr[\Delta c_{k+1} = \beta_{k+1} | \Delta m = \beta_0] = \sum_{\beta_k=0}^{max(\mathcal{M})} Pr[\Delta c_{k+1} = \beta_{k+1} | \Delta c_k = \beta_k] Pr[\Delta c_k = \beta_k | \Delta m = \beta_0]$$

As the characteristics form a Markov chain by Theorem 43 and we want to know when $\Delta c_{k+1} = \beta_{k+1}$ holds for any value preceding it. If we expand out our expression for $Pr[\Delta c_k = \beta_k | \Delta m = \beta_0]$ using our inductive hypothesis then we have

$$\begin{aligned} Pr[\Delta c_{k+1} = \beta_{k+1} | \Delta m = \beta_0] &= \sum_{\beta_k=0}^{max(\mathcal{M})} Pr[\Delta c_{k+1} = \beta_{k+1} | \Delta c_k = \beta_k] \sum_{j=1}^{k-1} \sum_{\beta_j=0}^{max(\mathcal{M})} \prod_{i=1}^k Pr[\Delta c_1 = \beta_i | \Delta m = \beta_{i-1}] \\ &= \sum_{j=1}^k \sum_{\beta_j=0}^{max(\mathcal{M})} Pr[\Delta c_{k+1} = \beta_{k+1} | \Delta c_k = \beta_k] \prod_{i=1}^k Pr[\Delta c_1 = \beta_i | \Delta m = \beta_{i-1}] \\ &= \sum_{j=1}^k \sum_{\beta_j=0}^{max(\mathcal{M})} \prod_{i=1}^{k+1} Pr[\Delta c_1 = \beta_i | \Delta m = \beta_{i-1}] \end{aligned}$$

With the last equality holding as any list of characteristics is a homogeneous Markov chain by Theorem 43. Therefore by induction this calculation holds. \square

However, while we can calculate these s-round characteristics quite easily in certain circumstances, it is often far too computationally intensive to calculate anything but lower bounds for s differential. We tend to calculate this lower bound by considering a number of good s-round characteristics. In the next section, we will be looking at how to find these s-round characteristics and to attain a good estimate for the lower bound of a differential.

5 Preparing for an Attack

5.1 Finding Good Differentials

Up until this point, the attack has been fairly deterministic. Given a Markov cipher we can calculate s -round characteristics using the definition of the sbox. Given some good characteristics we can calculate a lower bound for a differential and given that our lower bound holds with non-negligible probability we can, as we explain in Section 6, fairly reliably recover some key bits.

However, we do not have a process that allows us to generate the differentials that hold with the highest probabilities or even necessarily a process that gives us differentials with non-negligible probabilities. To discover these involves studying the structure of the cipher itself and therefore it is quite difficult to generate any successful general strategies beyond analysing the individual components of the cipher as we have done here and trying to use this information to calculate good differentials.

With that in mind, in this section we are going to outline some analysis performed on Unsub A to find good differentials. While the success of these strategies on Unsub A are not indicative of the success these will have on other ciphers, they can provide an example of the types of strategies one could use.

The logical starting point for analysing the cipher is to look at just a single round to try find good differentials. The advantage of this method is that it is often quite computationally easy to analyse just the one substitution and permutation layer in isolation. The disadvantage because you are analysing just a round in isolation, you have no guarantee that it is going to be a good differential for the entire cipher.

Our first line of investigation in this regard is looking at the likely single round characteristics. When calculating the probability that single round characteristics would hold, we constructed a difference distribution table D which held all the possible input differences and the corresponding frequency of each output difference. We can also use this table to identify the input difference to an sbox which will lead to an output difference from that sbox with the highest probability by finding the cell $D_{\Delta a, \Delta b}$ with the highest value. In other words, if the row Δa 's column Δb has the highest input in the entire of D , then the 8-bit difference that will most reliably be sent to another value is Δa and the value it is sent to is Δb .

Since sboxes generally have a smaller bit length than plaintext (8-bits for sboxes and 32-bits for the block for UnsubA), to get an 1-round characteristic from Δa and Δb above, we have to lengthen these differences somehow. One technique is to take the input difference $\Delta a'$ to be a pre-appended with 0's. Since the action of each sbox on its input bits is independent and an input difference 0 will be sent to 0 with probability 1 as discussed in our section on the substitution layer, this will give us that $\Delta a' \xrightarrow{S} \Delta b'$, where $\Delta b'$ is Δb preappended with 0's, will hold with the same probability that Δa lead to Δb given by

the difference distribution table D .

Another tactic relying on the behaviour of a single round is to choose input differences whose bits stay in one sbox after permutation. Let us say each sbox acts on s bits. Given a value a the i th sbox operates on the values $a_{is}, a_{is+1}, \dots, a_{(i+1)s-1}$. We say that a bit a_j with $qs \leq j < (q+1)s$ stays within its own sbox after permutation if the permuted bit $a_{\sigma(j)}$ has $rs \leq \sigma(j) < (r+1)s$ for some r smaller than the number of sboxes in the cipher. We can then examine all input differences that contain only bits that stay within an sbox. The idea behind this is that we are minimising the the number of active sboxes by just considering values that stay within one.

Example 48. To see these values for UnsubA, one examines the permutation network as detailed in Table 2a or illustrated in Figure 2. As both of these show, bits stay within their own sbox if they are the first or fifth bit from the left coming from that sbox. This means that the 8-bit numbers that stay within their sbox are 0b00000000, 0b00000001, 0b10000000, 0b10001000. We then can make 4^4 32-bit numbers by selecting four of these numbers, allowing a number to be selected multiple times, and concatenating them together. For example, both 0b00000000100010001000100000000001 and 0b10001000100010001000100010001000 are values that stay within their own sboxes.

Both of these methods only generate 1-round characteristics that hold with some probability. We still need a way to determine a likely $(r-1)$ -round differential with that input difference. To do this, we can construct an $(r-1)$ -round characteristic where each step of the characteristic is chosen based on the most likely one round characteristic given that input. Below is an outline an algorithm to do this.

Theorem 49 (Naive Maximum Characteristic). *Let us say we have some initial difference Δm . Let also assume that each sbox takes s bits and there are n sboxes. To generate a $(r-1)$ -round characteristic $\Delta m_0 = \alpha_0 \xrightarrow{\mathcal{R}} \dots \xrightarrow{\mathcal{R}} \alpha_{r-1}$, we follow the steps below:*

1. Set $i = 0$ and $\alpha_0 = \Delta m$
2. For every integer i smaller or equal to $r-1$ and larger than 0, we do the following:
 - (a) Split α_i into pieces each s bits long. Label these $\alpha_i^1, \alpha_i^2, \dots, \alpha_i^n$.
 - (b) Set α_{i+1}^j to the value that most likely follows from an input difference of α_i^j as can be looked up from the difference distribution table D by finding the largest column α_{i+1}^j in the row α_i^j of D .
 - (c) Set α'_{i+1} to the concatenation of $\alpha_{i+1}^1, \dots, \alpha_{i+1}^n$
 - (d) Set α_{i+1} to $P(\alpha'_{i+1})$ where P is the permutation function of our SPN block cipher.
3. Your $(r-1)$ round characteristic is the $\alpha_0, \alpha_1, \dots, \alpha_{r-1}$ produced by this algorithm.

We note that this process might not generate the best possible characteristic - suppose we have some characteristic $a \xrightarrow{\mathcal{R}} b \xrightarrow{\mathcal{R}} c$ generated using the above algorithm. There is nothing to say that we cannot find a b' such that $a \xrightarrow{\mathcal{R}} b'$ holds with lower probability than $a \xrightarrow{\mathcal{R}} b$ but $b' \xrightarrow{\mathcal{R}} c$ is much higher. However, with a decent initial difference the algorithm above can create a difference good enough to break our toy cipher. If necessary, one could iteratively improve such characteristics by trying to find values such as b' that improve the whole characteristic, possibly changing the differential in the process. Unfortunately, to do this effectively requires potentially quite a large amount of computer power and is generally infeasible.

Example 50. For UnsubA, we can calculate a difference distribution table using the sbox outlined in Table 1a. If we do this, we see that we get that the input difference 0x20 goes to 0x2 with probability ≈ 0.0859375 . If we take 0x00000020 as our initial difference and take each step in the characteristic as the one that follows from the previous with the largest probability as outlined in Theorem 49 then you get the following 5-round characteristic.

$$\begin{aligned} 0x20 &\xrightarrow{S} 0x2 \xrightarrow{P} 0x100 \xrightarrow{S} 0xc000 \xrightarrow{P} 0x8080000 \xrightarrow{S} 0xbebe0000 \\ &\xrightarrow{P} 0xf050f0a0 \xrightarrow{S} 0xd005d02c \xrightarrow{P} 0x89990298 \xrightarrow{S} 0x6e8d4217 \\ &\xrightarrow{P} 0x70d9c513 \xrightarrow{S} 0x182eee01 \xrightarrow{P} 0x5c1c3c81 \end{aligned}$$

Here we have removed the key addition step as we showed it does not impact the characteristic.

This characteristic holds with probability $1.13891793191 \times 10^{-23}$ which means that the differential (0x20, 0x5c1c3c81) holds with at least that probability. This is unfortunately far too low to use for a differential attack. So we will have to try another method to find a differential that will.

5.2 Filtering

When introducing differentials, we noted that it is difficult to calculate the actual probability a differential holds and instead we would use a couple good characteristics together to try find a good lower bound for this probability. With this probability known, we can now try partial decryption of pairs of ciphertexts, looking for last-round keys that lead to our predicted difference in the $(r - 1)$ th round. We can look at this as a signal to noise problem: we want to differentiate between the signal of the possible right pairs giving us potential key bits and the noise of the wrong pairs that do not give us any information. By picking a good differential, we boost this signal. Now we are going to discuss how we

can filter out wrong pairs before we partially decrypt them in order to reduce the noise.

Like picking good differentials, there is no deterministic way to construct good filtering. While filtering is important, it is very dependent on the differential chosen.

The method we used while analysing Unsub A is to examine what differences encrypted right pairs could have by examining the last difference in the characteristic. Below we sketch an algorithm for doing this.

Theorem 51 (Naive Filtering). *Given an $(r - 1)$ -round differential $(\Delta m, \Delta l)$ in a cipher with n sboxes each with s bits, we want to calculate all the possible Δc a right pair (c, c') could have. Let us define*

$$D_a = \{b \mid D_{a,b} \neq 0\}$$

Then, we can do the following:

1. Take Δl and divide it into n s -bit numbers $\Delta l_1, \Delta l_2, \dots, \Delta l_n$.
2. Calculate $c_i = D_{l_i}$ for $1 \leq i \leq n$.
3. Take the set C to be all numbers x such that the s_j through $s(j + 1)$ bits make up an element of c_j for $0 \leq j \leq n$.
4. We say that Δc follows from Δl if Δc is equal to some element of C after it has been permuted by the permutation layer of our cipher, i.e. Δc follows from Δl iff. $(\exists x \in C) \Delta c = P(x)$.

Put loosely: we find all the possible differences that Δl can make by examining all the possible differences its bits can make when being fed into sboxes and then finding all the permutations of these possibilities.

We then construct a function $F : \mathcal{M} \rightarrow \{0, 1\}$ such that

$$F(x) = \begin{cases} 1 & x \text{ follows from } \Delta l \\ 0 & \text{otherwise} \end{cases}$$

Example 52. Let us look at the 5-round differential (0x11000000,0x40004). To apply the filtering mentioned above, we need to find all the possible output differences 0x40004 can lead to. To do this, we first have to split it into 8-bit chunks and look each chunk up in the difference distribution table and find all the non-zero entries. If we do this, we get the following

Hex Value	Number of non-zero entries
0x0	1
0x4	99
0x0	1
0x4	99

As P is bijective, this means that from 0x40004 we can achieve $99^2 = 9801$ possible output differences. While this seems quite large, if we consider that there are $2^{31} = 2147483648$ possible output differences, this means that we can eliminate $2147483648 - 9801 = 2147473847$ output differences if seen, or we only need to examine $\approx 4.6 \times 10^{-4}\%$ of the output pairs that can be produced.

6 Attacking a Cipher using Differential Cryptanalysis

As we stated in Section 4, in this section we are going to step through how to attack a block cipher when you have identified a good differential and determined a filtering system as outlined in Section 5. We will then implement this attack on UnsubA with a randomly generated key to show that such an attack is practical.

6.1 Generic One Round Attack

The initial generic attack we present will simply be an improvement over the brute force technique: instead of having to guess every round key, we will attack the last one round key. This is useful for ciphers where the last round key is smaller than the full key for the cipher. We will then outline a simple extension to this attack in Section 6.1.3 which improves this by allowing us to attack bits of the last round key rather than the entire key.

6.1.1 Algorithm

Let us take an iterated block cipher with r rounds encrypting under a secret key k_r . Assuming that we have a $(r-1)$ -round differential $(\Delta m, \Delta l)$ which holds with probability p , we perform the following attack on the N unique pairs of tuples (m, c) and (m', c') where $m \oplus m' = \Delta m$ and $E(k_r, m) = c, E(k_r, m') = c'$. Also let us assume we have some filter as described in Section 5.2 defined as the function $F : \mathcal{M} \times \mathcal{M} \rightarrow \{0, 1\}$ where $F(c, c') = 1$ if they are possibly a right pair and 0 otherwise. If we do not have such a filter, assume that $F(x, x') = 1$ for all $(x, x') \in \mathcal{M} \times \mathcal{M}$. Let define a round of decryption of a ciphertext c under our key as $R^{-1}(c, k_r)$ with R being a function denote a round of the function. We are also going to need a counter for each key $t \in \mathcal{K}$.

With these, we have the following attack:

Theorem 53 (Generic Differential Attack on Iterated Block Cipher). *To perform the attack, we follow the algorithm below.*

1. *For each pair of ciphertexts (c, c') corresponding to the plaintexts (m, m') respectively, if $F(c, c') = 0$, remove them from the current set of pairs we are analysing.*
2. *For all the remaining pairs (c, c') and for each value $t \in \mathcal{K}$ we do the following*
 - (a) *Compute $\Delta l' = R^{-1}(c, t) \oplus R^{-1}(c', t)$*
 - (b) *If $\Delta l' = \Delta l$, increment counter T_t*
3. *Identify the counter T_i which has the expected correct value Np .*
4. *Guess that $k_r = i$.*

6.1.2 Explanation

Let us walk through this attack step by step, explaining why what we are doing gives us a decent guess for the key required.

In the first step, we are simply applying the filtering method to our pairs to remove wrong pairs. Wrong pairs by definition do not follow our differential and therefore cannot give the correct candidate key when partially decrypted. Using the analogy given in filtering: if we left the wrong pairs in, we simply increase the noise, making it harder to find the correct candidate key. As such, we want to remove as many wrong pairs as possible through this filtering.

Now, with characteristics, we said a plaintext pair was right or wrong depending on whether it followed that characteristic exactly or not. With respect to a differential and a key, a plaintext pair is called a right pair if when it is encrypted using the key, the difference predicted by the differential occurs. Therefore, if the pair is a right pair, then at least one of the candidate keys generated will be k_r as that key will give the expected difference that this pair follows. Conversely, a wrong pair cannot have k_r associated with it. If it did, then using the pairs defined above, the difference between the partially decrypted values of (l, l') under k_r would give Δl , meaning this wrong pair was actually a right pair. This means that the candidate keys generated by a wrong pair cannot include k_r .

The next step in our process is to compute the decryption of each ciphertext with each candidate key and compare it to the predicted difference from the differential. If this value is the same as the one predicted by the differential, then it is possibly a right pair and so we record its value in a counter. If it does not equal the expected difference, then it is a wrong pair and so we just discard it.

The most important step of this attack is identifying the guess for our key. This is done by choosing the counter with the predicted distribution of pairs from the differential. Given a pair of plaintexts (m, m') and corresponding ciphertexts (c, c') with $m \oplus m' = \Delta m$, let us call the keys t that give $\Delta l = R^{-1}(c, t) \oplus R^{-1}(c', t)$ the candidate keys generated by (m, m') as for these keys, the partial decryption of (c, c') looks like a right pair.

We then assume what is known as *wrong-key randomisation*: that the set of candidate keys for a wrong pair are independent of k_r and generated randomly. Assuming that we have a differential that holds with a significantly large probability, under wrong-key randomisation we assume the wrong keys will all be suggested roughly the same number of times with k_r suggested a disproportionately large number of times, allowing us to easily differentiate k_r from the wrong keys.

To identify this frequency, we use the probability associated with our differential. We have that our $r - 1$ round differential holds with a probability p , therefore we expect that for N pairs with the input difference Δm that we will receive pN pairs that have a difference of Δl after $r - 1$ rounds, i.e. pN right pairs. Therefore we expect that if our key candidate key is correct that it will partially decrypt pN values to Δl . Therefore the counter with

pN values, or the one with the closest value to it, is our best guess for a key.

This means that if we follow our generic attack with a good differential, we should be able to find out correct key k_r .

6.1.3 Extensions of Generic Attack

This attack is often specialised by modifying the second step to instead look for partial keys rather than full keys. This means instead of looking for all the bits of the key, we instead just search for some bits of it. This approach is often viable when looking at cipher like unsubA where the key addition is done with a bit wise operation such as bit wise xor. In this approach, instead of performing an inverse round on the entire value, we perform it only on the bits of the key we are searching for - this means that we often end up looking for partial keys the size of the sboxes to make this easier. This is often performed in conjunction with $(r - 1)$ -round characteristics which only have non-zero bits in the expected last difference in one sbox. You then can target that sbox as the range from which to recover key bits. We then compare the resulting bits with the corresponding bits of the differential, incrementing counters as per usual. We then repeat this attack until we have all the keybits we want, up to the entire last round key.

For a cipher like UnsubA, due to the easily reversible nature of the key scheduling algorithm, we just need to discover one key in order to get all the keys for the round. For a block cipher where the key size is bigger than the block size, we might want to uncover more than one rounds worth of keys. This generic attack can extended to do this quite easily: instead of looking at a $r - 1$ round differential, you instead look at $r - i - 1$ where i are the number of round keys you already have. To do this, you will need to recover i full round keys through either the generic attack of recovering all of the bits through repeated applications of the extended attack covered above. Then, before decrypting a ciphertext with the candidate key, you decrypt the ciphertext with the available i round keys. With this, you can discover the key for that round, repeating the process again to find another as necessary.

6.2 Attack on Unsub A

We are going to use our extended attack above to recover 16 key bits from a reduced version of UnsubA - a version which instead of having 6 rounds has 4. The 3-round differential we are going to use for this attack is $(0x11000000, 0x20200000) = (\Delta m, \Delta l)$. We obtained this differential by considering the 3-round characteristic generated by combining our two methods in Section 5: we took all the possible values that stay within an sbox and generated 3-round characteristics by chaining together the largest 1-round characteristic for the given input difference as detailed in Theorem 49. The result gave us the characteristic

$$0x11000000 \xrightarrow{\mathcal{R}} 0x00000040 \xrightarrow{\mathcal{R}} 0x00010000 \xrightarrow{\mathcal{R}} 0x20200000$$

which holds with probability $5.14984130859375 \times 10^{-5}$.

For the attack below, we are going to label our bits from the right to the left starting from 0. So the number 0b101 has 1 for its 0th bit, 1 for its 1th bit and 0 for its 3rd bit. We are also going to label our sboxes 0 through 4 starting with the one on the right as 0.

Let us call the key we are going to try recover information from, the fifth key added to our plaintext during encryption, as k . Since our expected output for the third round is zero for the first 16 bits, we cannot recover all of this key. However, we can try and recover 16 bits of k by using the non-zero last half of our expected difference.

Trying to recover specific bits makes this attack slightly more complex than our first attack but not by much. First, we have to pick an sbox range to attack. Our expected output dictates we choose either sbox 2 or 3 so let us say our choice is i with $i = 2$ or $i = 3$. Now, given that we know we want to look at the $8i$ to $8(i + 1) - 1$ bits of Δl , how do we know which key bits we want to recover? Well, we can simply take each of the bit locations in this range, permute them with the permutation σ as defined by the bit permutation in Table 2a for UnsubA's permutation layer. Then we have a list of bits in the key that we can recover using the values that correspond with our sbox i .

Using this information, we can make some changes to our generic algorithm to make guessing the secret key easier. To do this, note that $\mathcal{R}^{-1}(c, k)$ for secret key k and ciphertext c is equivalent to saying $S^{-1}(P^{-1}(k \oplus c))$ which equals $S^{-1}(P^{-1}(c) \oplus P^{-1}(k))$. Now, since we are going to iterate through all the values of k and since P is a bijection, this is equivalent to iterating through all the values t in the expression $S^{-1}(P^{-1}(c) \oplus t)$.

Now, let us restrict t to only having non-zero bits for bits in the range $8i$ to $8(i + 1) - 1$. Since sboxes work on blocks of 8 bits, this means that $S^{-1}(P^{-1}(c) \oplus t)$ will be the partial decryption of c under a key with t 's bits in its non-zero range for the i th sbox. But, this is the only range we care about as this is the range we are going to compare to our expected output difference in order to recover a key. This means that we can simply disregard any bits of $S^{-1}(P^{-1}(c) \oplus t)$ if they are not in the range $8i$ to $8(i + 1) - 1$ as they are not touched by our key. If we do this for all $2^8 - 1$ possible t values, and compare them against a similarly truncated version of Δl , we are performing a differential attack for 8 key bits that correspond with the i th sbox. The locations of the bits we recover from k will be the all the numbers in the range $8i$ to $8(i + 1) - 1$ permuted by σ .

With all this in mind, we have the following algorithm in order to recover our key bits. In our attacks, we used no filtering as our differential was good enough on its own. Please note that this is simply the generic attack with the modifications as specified above so we will assume we have all the conditions as specified in Section 6.1.1 and that we are searching for bits of the key $k_r \in \mathcal{K}$.

To perform the attack on UnsubA, we follow the algorithm below.

1. For all the pairs (c, c') and for each value $t \in \mathcal{K}$ where t only has non-zero bits in the range $8i$ to $8(i+1) - 1$, we do the following
 - (a) Compute $\Delta' = S^{-1}(P^{-1}(c) \oplus (t)) \oplus S^{-1}(P^{-1}(c') \oplus (t))$.
 - (b) If the bits from $8i$ to $8(i+1) - 1$ for $\Delta' = \Delta l$ are equal, increment counter T_t
2. Identify the counter T_i which has the largest number of entries.
3. Guess that j th bit of i is the $\sigma(j)$ th bit of k_r for all j in the range $8i$ to $8(i+1)$.

In this project, we ran the this algorithm with 10000 pairs for the 2nd and 3rd sboxes by simply setting i to 2 and 3. The result of this was that we are able to successfully recover 16 key bits of the secret key in our four round reduced cipher. The attacks each took just over a minute on a Intel Core i3 Processor.

6.2.1 Data Complexity and chance for Success

For a differential attack, we need two conditions to hold:

1. There must be a reasonable chance that the differential holds for at least one pair.
2. The correct value of k_r must be suggested at a significantly different rate than other values.

If the first condition does not hold then all the keys suggested will be incorrect and therefore the attack cannot yield the correct key. To ensure that this is the case for a differential that holds with probability p , we will then need to choose about p^{-1} pairs. In fact, we can use a theorem by Selcuk [5] to gauge how successful an attack will be for a given number of pairs. Before we do that though, we need to look at the fulfilling of the second criteria: i.e. trying to predict at what rate the correct key is suggested compared to other keys.

Definition 54. *Let us consider one iteration of the differential attack, i.e. doing the attack for one pair. If p is the probability that the correct key is suggested from this pair and p_r is the average probability that an incorrect key is suggested, then **signal-to-noise ratio** is*

$$S_N = \frac{p}{p_r}$$

As p here will be the probability that our differential holds, and p_r being the probability we produce a key we have to eventually ignore we see that this definition makes sense as p

is the probability we will produce a “signal” every iteration while p_r is the probability we will produce noise.

Let us consider how this relates to the condition we are examining.

$S_N \approx 1$ Under wrong-key randomisation we will have that the correct key is produced the same number of times as the other keys. This means that differentiating it from the other keys will be difficult or impossible.

$S_N \gg 1$ Here, the correct key should be produced far more often than the incorrect keys which means that it the correct key should be the one suggested more often than the others.

Now, calculating S_N is easy if we do not use filtering or if we are attacking the whole key. However, since both of these allow us to greatly increase our chance of success, it is worth analysing how to calculate S_N under these extended attacks.

Theorem 55. *Let λ be the probability that a randomly chosen pair survives filtering, γ be the average number of key values suggested by each pair that survives filtering and κ denote the number of bits of the key we are trying to recover. Then for a differential that holds with probability p we have*

$$S_N = \frac{p(2^\kappa - 1)}{\lambda\gamma - p}$$

Proof. We have defined

$$S_N = \frac{p}{p_r}$$

where p_r is the average probability that an incorrect key is suggested. To calculate S_N is therefore to derive an expression for p_r .

Now, firstly let us note that $\lambda \geq p$ as filtering can only remove wrong pairs and the chance that a pair is a right pair is, of course, p . They are equal if filtering removes all wrong pairs but otherwise filtering will leave some wrong pairs which means the chance a pair survives filtering must be higher than the chance it is a right pair. This also means that the probability that any given wrong pair survives filtering is $\lambda - p$.

Before trying to calculate p_r , let us look at the number of incorrect keys generated on average by a pair. That pair has a $\lambda - p$ chance of being a wrong pair and it will generate γ keys, on average. As discussed in the attack, all γ of these keys will be wrong. Similarly, the chance of a pair being right is p and the number of wrong keys that pair will generate on average is $\gamma - 1$, i.e. the number of keys it will generate on average less the correct one. This means that the number of incorrect keys generated on average by a randomly selected pair is

$$(\lambda - p)\gamma + p(\gamma - 1)$$

Since p_r is the chance that a given key that was suggested is wrong, and we have assumed that incorrect key values are uniformly distributed over all possibly $2^\kappa - 1$ incorrect values of the key through wrong-key randomisation, we have that

$$p_r = \frac{(\lambda - p)\gamma + p(\gamma - 1)}{2^\kappa - 1}$$

Substituting this into our expression for S_N gives us

$$S_N = \frac{p}{p_r} = \frac{p(2^\kappa - 1)}{(\lambda - p)\gamma + p(\gamma - 1)}$$

as desired. □

It is worth noting that even with a good S_N , a differential attack does not necessarily succeed. This is because the Hypothesis of Stochastic Key Equivalence and wrong-key randomisation do not necessarily hold. Also, if the round keys are not chosen uniformly at random as desired, our results are not valid as this was assumed at the start of our analysis. Though if these do hold, and provided the attacker can find a good differential, these form quite a powerful way of breaking the confidentiality of ciphers.

With this definition of S_N , we can now express the chance that a differential attack will succeed. As the proof for this would be a project within itself, the result will simply quoted from Selccuk's paper [5]. Before we can cite that result, we need to explain some of the concepts behind it.

Definition 56 (Advantage). *If a differential attack on an b bit key gets the value ranked among the top r out of 2^b candidates then we say that the attack obtained a $(b - \log r)$ -bit **advantage** over exhaustive search.*

Saying that your differential attack has the correct key as the top value is the same as saying that the attack gained a b bit advantage over a b bit key as $\log 1 = 0$.

Definition 57 (Normal distribution function). *The normal distribution function $\Phi(x)$ is defined as*

$$\Phi(x) := \frac{1}{\sqrt{2\pi}} \int_{-x}^x e^{-t^2/2} dt$$

Its inverse is

$$\Phi(x)^{-1} = \sqrt{2} \operatorname{erf}^{-1}(2x - 1), \operatorname{erf}^{-1}(x) := \sum_{k=0}^{\infty} \frac{c_k}{2k+1} \left(\frac{\sqrt{\pi}}{2} x \right)^{2k+1}$$

Theorem 58 (Selccuk). *Let P_S be the probability that a differential attack on an m -bit key, with a characteristic of probability p and a signal-to-noise ratio S_N , and with N plaintext-ciphertext pairs delivers an a -bit or higher advantage. Assuming that the key counters are*

independent and that wrong-key randomisation holds, we have for sufficiently large m and N

$$P_S = \Phi \left(\frac{\sqrt{pNS_N} - \Phi^{-1}(1 - 2^{-a})}{\sqrt{S_N + 1}} \right)$$

A Unsub A Construction

A.1 Cipher Description

In this project, we analyse a toy cipher⁴ created by Riaal Domingues [2]. Its block and key size are both 32-bits long and it uses a Substitution Permutation Network - its encryption is accomplished through a series of non-linear substitutions and linear permutations on the bits of the input. For some visualisations of UnsubA, see Figures 1, 2 and 3. These are all designed by myself and the source code for generating these diagrams will be released under a BSD licence at <http://github.com/steveryb>.

A.2 Round Structure

This cipher consists of 6 rounds with each round consisting of a key addition, substitution and permutation in that order, with the output of the one step forming the input for the next. After these 6 rounds, there is a post whitening key added at the end using the same key addition process.

A.2.1 Key Addition

Let the input for round i be A^i with i th key for the round being K^i . The key K^i is generated from a single secret key K using the key expansion algorithm detailed in Section A.2.4. The output B^i after key addition is $A^i \oplus K^i$ with \oplus being binary xor as defined in Section 23.

A.2.2 Substitution

Given a B^i as input after key addition, the substitution step takes the $B^i = b_0^i b_1^i \dots b_{30}^i b_{31}^i$ where b_j^i is the j th bit of B^i when counting from the left and splits it into four equal length inputs $B_{0,7}^i, B_{8,15}^i, B_{16,23}^i, B_{24,31}^i$ where $B_{l,m}^i = b_l^i b_{l+1}^i \dots b_{m-1}^i b_m^i$. Each of these groups of $B_{l,m}^i$ is then substituted according to the invertible sbox S as defined in Table 1a to create $C_{l,m}^i$. This is to say that $S(B_{l,m}^i) = C_{l,m}^i$. The output of this round is C^i which is defined as $C^i = c_0^i c_1^i \dots c_{30}^i c_{31}^i = C_{0,7}^i C_{8,15}^i C_{16,23}^i C_{24,31}^i$.

A.2.3 Permutation

Given C^i as given after the substitution step, we permute the bits of C^i to get D^i according to the look up table in Table 2a. The way we do this is to define an automorphism P on

⁴It is a toy cipher in the sense that it was created so that cryptographers could try break it for practice - it was not made to be secure in other words.

$X = \{0, 1, 2, \dots, 30, 31\}$ by the the permutation look up table. Then, we have for every bit d_j in the representation of D^i that $d_j = c_{P(j)}$.

A.2.4 Key Schedule

While this is not explicitly part of the round structure, it is needed to generate the keys required by the Key Addition step. When encrypting a data, a secret key K is given. To generate the key K^i required for Round i , we define a permutation on the bits of $K = k_0k_1 \dots k_{30}k_{31}$. For $f : \{0, 1, \dots, 30, 31\} \times \{0, 1, 2, \dots, 6\} \rightarrow \{0, 1, \dots, 30, 31\}$, we define

$$f(b, i) = b + i \bmod 32$$

With this, we have that $K^i = k_{f(0,i)}k_{f(1,i)} \dots k_{f(30,i)}k_{f(31,i)}$. This permutation is taking bit shifts of the keys of k , essentially “rotating” the bits.

A.3 Encryption and Decryption Algorithms

To encrypt a value m with a key k , you perform 6 rounds as explained above, starting at Round $i = 0$ with $A = A^0$ and $k = K^0$.

1. Take A^i , perform the key addition with K^i to get B^i .
2. Take B^i , take it through the substitution boxes (sboxes) to get C^i .
3. Take C^i and permute it according to the permutation layer to get D^i .
4. If $i \neq 5$: set $A^{i+1} = D^i$ and perform this process again for round $i + 1$ from Step 1. Otherwise, stop.

As a final step, we perform one last key addition to D^5 using K^6 to get the final c value, i.e. $E_{UnsubA}(m, k) = c$.

To decrypt values, since all these processes are invertible, we simply run this process in reverse:

Given a value c , we perform a key addition to it using K^6 to get D^5 . Then, starting at Round $i = 5$, we perform the following:

1. Take D^i and permute it according to the permutation network defined in Table 2b to get C^i
2. Take C^i , take it through the inverse substitution boxes defined in Table 1b to get B^i

3. Take B^i , perform the key addition with K^i to get A^i
4. If $i \neq 0$: set $D^{i-1} = A^i$ and perform this process again for round $i - 1$ from Step 1. Otherwise, stop.

We then get back A^0 which we call m , the plaintext of c . In other words, $D_{UnsubA}(c, k) = m$ with m being found as the process above stipulates.

References

- [1] Eli Biham and Adi Shamir. Differential cryptanalysis of des-like cryptosystems. *Journal of CRYPTOLOGY*, 4(1):3–72, 1991.
- [2] Riaal Domingues. Unsub: Toy ciphers for games in cryptanalysis. *Self published*, 2011.
- [3] Charles M. Grinstead and J. Laurie Snell. *Grinstead and Snell’s Introduction to Probability*. American Mathematical Society, 2006.
- [4] L.R. Knudsen and M.J.B. Robshaw. *The Block Cipher Companion*, pages 109–126,145–165. Springer-Verlag Berlin Heidelberg, 2011.
- [5] Ali Aydın Selçuk. On probability of success in linear and differential cryptanalysis. *Journal of Cryptology*, 21(1):131–147, 2008.
- [6] Claude E Shannon. Communication theory of secrecy systems. *Bell system technical journal*, 28(4):656–715, 1949.
- [7] James L. Massey Xuejia Lai and Sean Murphy. Markov ciphers and differential cryptanalysis. Institute for Signal and Information Processing, Swiss Federal Institute of Technology, CH-8092 Zurich, Switzerland, 1991.

x	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
0	40	80	b8	b7	3f	d9	24	da	5f	ff	60	9f	81	a0	db	0a
1	3d	41	3e	7f	25	d8	5e	dd	23	9e	61	dc	82	0b	a1	83
2	b9	42	3c	d7	00	5d	26	de	7e	9d	e0	0c	62	a2	22	84
3	3b	43	b6	5c	ba	7d	df	bb	27	9c	fe	e1	63	85	a3	21
4	3a	bc	e4	44	d6	01	5b	7c	28	9b	0d	e2	20	a4	e3	64
5	45	39	a6	d5	bd	e5	a5	5a	7b	e6	0e	9a	1f	d4	65	86
6	66	b5	38	46	59	99	47	02	c2	0f	e7	1e	c0	be	e8	7a
7	b4	fd	67	37	a7	98	10	48	09	c1	1d	e9	79	bf	b3	87
8	68	58	97	a8	69	11	fc	36	49	03	29	ea	a9	78	fb	88
9	96	eb	57	12	6a	ab	35	c4	1c	fa	4a	c3	2a	77	aa	89
a	6b	ac	95	ec	56	c6	76	34	ed	f9	04	c5	ad	4b	2b	8a
b	c9	6c	c7	94	13	33	75	d2	ee	55	ef	4c	ae	6d	b2	2c
c	c8	d1	14	4d	32	74	93	1b	f8	54	d3	05	6e	f0	2d	8b
d	73	15	31	4e	d0	1a	92	53	ca	f7	6f	f2	8c	06	f1	2e
e	72	f3	30	4f	16	f4	cf	91	52	cb	19	70	8d	ce	07	8e
f	2f	50	af	cc	b0	17	f5	51	cd	18	f6	90	b1	71	8f	08

(a) Sbox

x	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
0	24	45	67	89	aa	cb	dd	ee	ff	78	0f	1d	2b	4a	5a	69
1	76	85	93	b4	c2	d1	e4	f5	f9	ea	d5	c7	98	7a	6b	5c
2	4c	3f	2e	18	06	14	26	38	48	8a	9c	ae	bf	ce	df	f0
3	e2	d2	c4	b5	a7	96	87	73	62	51	40	30	22	10	12	04
4	00	11	21	31	43	50	63	66	77	88	9a	ad	bb	c3	d3	e3
5	f1	f7	e8	d7	c9	b9	a4	92	81	64	57	46	33	25	16	08
6	0a	1a	2c	3c	4f	5e	60	72	80	84	94	a0	b1	bd	cc	da
7	eb	fd	e0	d0	c5	b6	a6	9d	8d	7c	6f	58	47	35	28	13
8	01	0c	1c	1f	2f	3d	5f	7f	8f	9f	af	cf	dc	ec	ef	fe
9	fb	e7	d6	c6	b3	a2	90	82	75	65	5b	49	39	29	19	0b
a	0d	1e	2d	3e	4d	56	52	74	83	8c	9e	95	a1	ac	bc	f2
b	f4	fc	be	7e	70	61	32	03	02	20	34	37	41	54	6d	7d
c	6c	79	68	9b	97	ab	a5	b2	c0	b0	d8	e9	f3	f8	ed	e6
d	d4	c1	b7	ca	5d	53	44	23	15	05	07	0e	1b	17	27	36
e	2a	3b	4b	4e	42	55	59	6a	6e	7b	8b	91	a3	a8	b8	ba
f	cd	de	db	e1	e5	f6	fa	d9	c8	a9	99	8e	86	71	3a	09

(b) Inverse Sbox

Table 1: A lookup table defining the operation of a substitution box (sbox) and its inverse as implemented in Unsub A. Note all values are represented in hexadecimal. Substituting a number A with bits $a_0a_1a_2a_3a_4a_5a_6, a_7$ with this sbox will result in the number in the cell in the $a_0a_1a_2a_3$ th row and the $a_4a_5a_6a_7$ th column. For examples of substitution, see Table 3

x	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
0	00	08	16	24	01	09	17	25	02	10	18	26	03	11	19	27
1	04	12	20	28	05	13	21	29	06	14	22	30	07	15	23	31

(a) Permutation

x	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
0	00	04	08	12	16	20	24	28	01	05	09	13	17	21	25	29
1	02	06	10	14	18	22	26	30	03	07	11	15	19	23	27	31

(b) Inverse Permutation

Table 2: A lookup table defining the operation of the permutation network as implemented in Unsub A. Note all values are represented in hexadecimal. Permuting a number A with bits $a_0a_1a_2a_3a_4$ with this permutation will result in the number in the cell in the a_0 th row and the $a_1a_2a_3a_4$ th column. Some examples of this are in Table 4

Value	Substituted Value
0x1f	0x83
0x3e	0xa3
0xc3	0x4d
0x8d	0x78
0x2d	0xa2
0x98	0x1c
0xe7	0x91
0x11	0x41
0x58	0x7b
0xbb	0x4c

(a) Substituted Values

Value	Inverse-Substituted Value
0xad	0xac
0xa9	0x8c
0xcc	0xf3
0x6e	0xcc
0x5a	0x57
0xb0	0xf4
0x83	0x1f
0x4a	0x9a
0x88	0x8f
0x9a	0x5b

(b) Inverse Substituted Values

Table 3: Some example values for substitution and inverse substitution layer of UnsubA

Value	Permuted Value
0x81e2e4d8	0xab2e3842
0x1a929fa3	0x6e0457ad
0xcb80f433	0xe88c4b4b
0xd861a67e	0xc9a72f92
0xa33cb564	0x9817ea6c
0x9aa7c1a9	0xeb187295
0xca8f4ff4	0xf69f5616
0x2c28dab3	0x5e48a70b
0x5c58a590	0x5ae408a6
0x3b890d79	0x7506c2d7

(a) Permuted Values

Value	Inverse-Permuted Value
0x77fd6cbc	0x5efd7f8c
0x53cfe47b	0x6f3956dd
0xe625e87a	0xabf13c94
0xd0d0a779	0xed3d1223
0x342a0dc1	0x11c86a43
0x32b8c0ef	0x73dc5191
0x21de71d2	0x57a7445a
0x1c9205aa	0x501c9a52
0x99b47f05	0xc26ea72b
0x1a91d961	0x631ea087

(b) Inverse-Permuted Values

Table 4: Some example values for permutation and inverse permutation layer of UnsubA

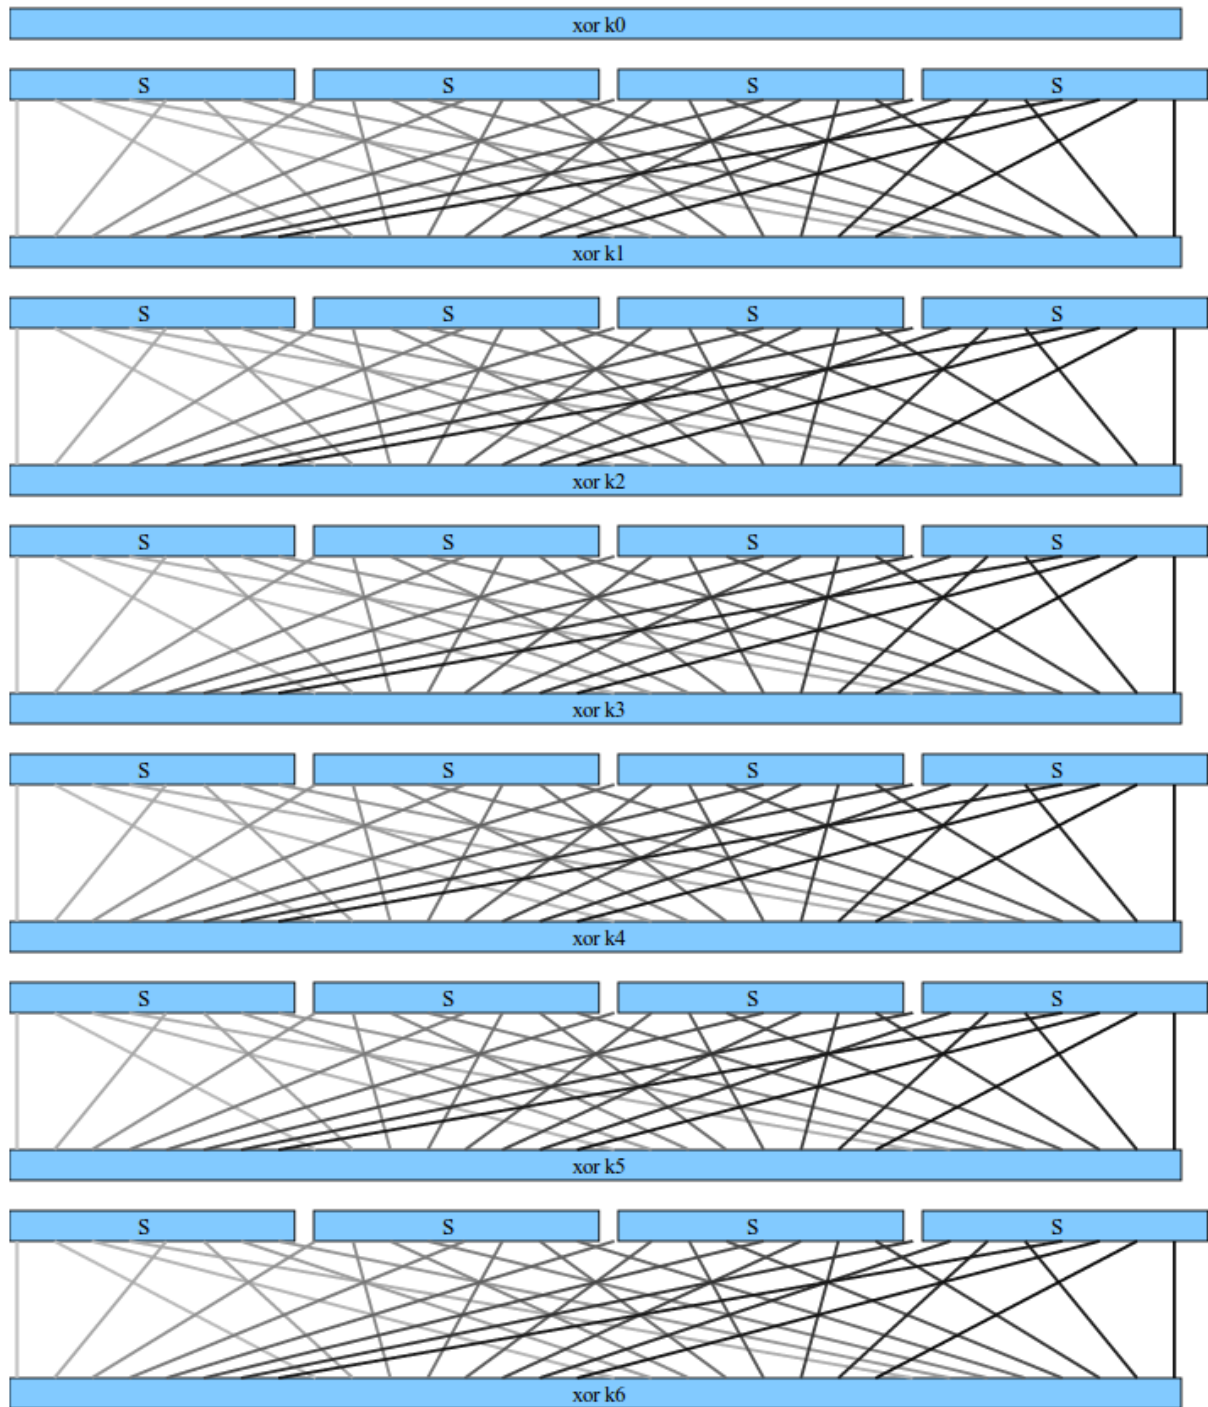


Figure 1: A visualisation of UnsubA showing how the bits get distributed under the permutation layer.




Figure 2: A visualisation of UnsubA illustrating where the outputs of each sbox is distributed by the permutation layer. Each different colour indicates a different sbox origin, the lines indicate where each bit of the output of the sbox is permuted to.

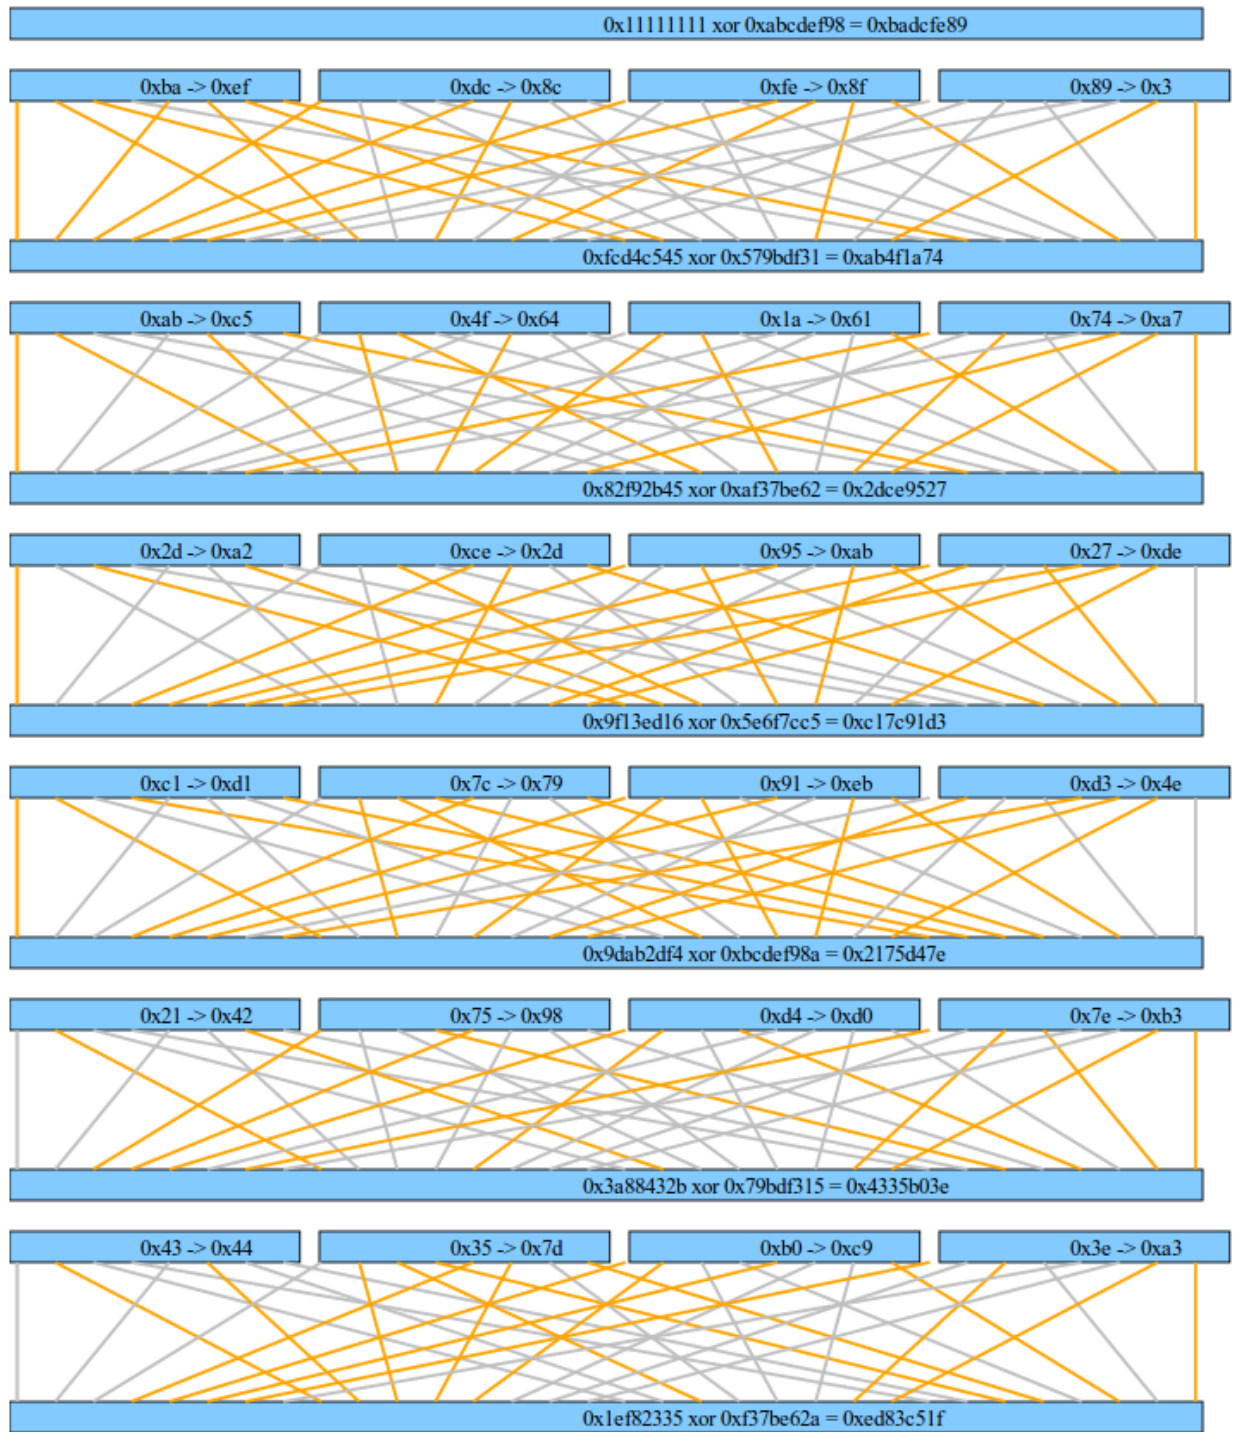


Figure 3: A visualisation of UnsubA detailing the encryption of 0x11111111 under the key 0xabcdef98. The long horizontal boxes indicate the key addition stages, the boxes below them indicate the values given to the sboxes and the values they output. The coloured lines give the non-zero bits given to the permutation layer and where they get distributed to.